

# Scheduling and Analysis of Limited-Preemptive Malleable Gang Tasks

Joan Marcè i Igual

Geoffrey Nelissen

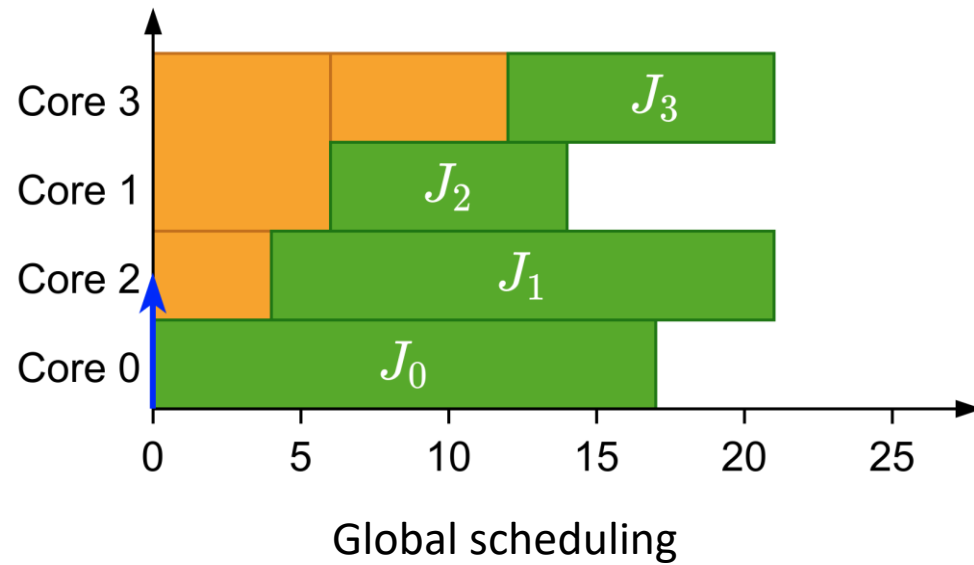
Mitra Nasri

Paris Panagiotou

24<sup>th</sup> of February, 2020

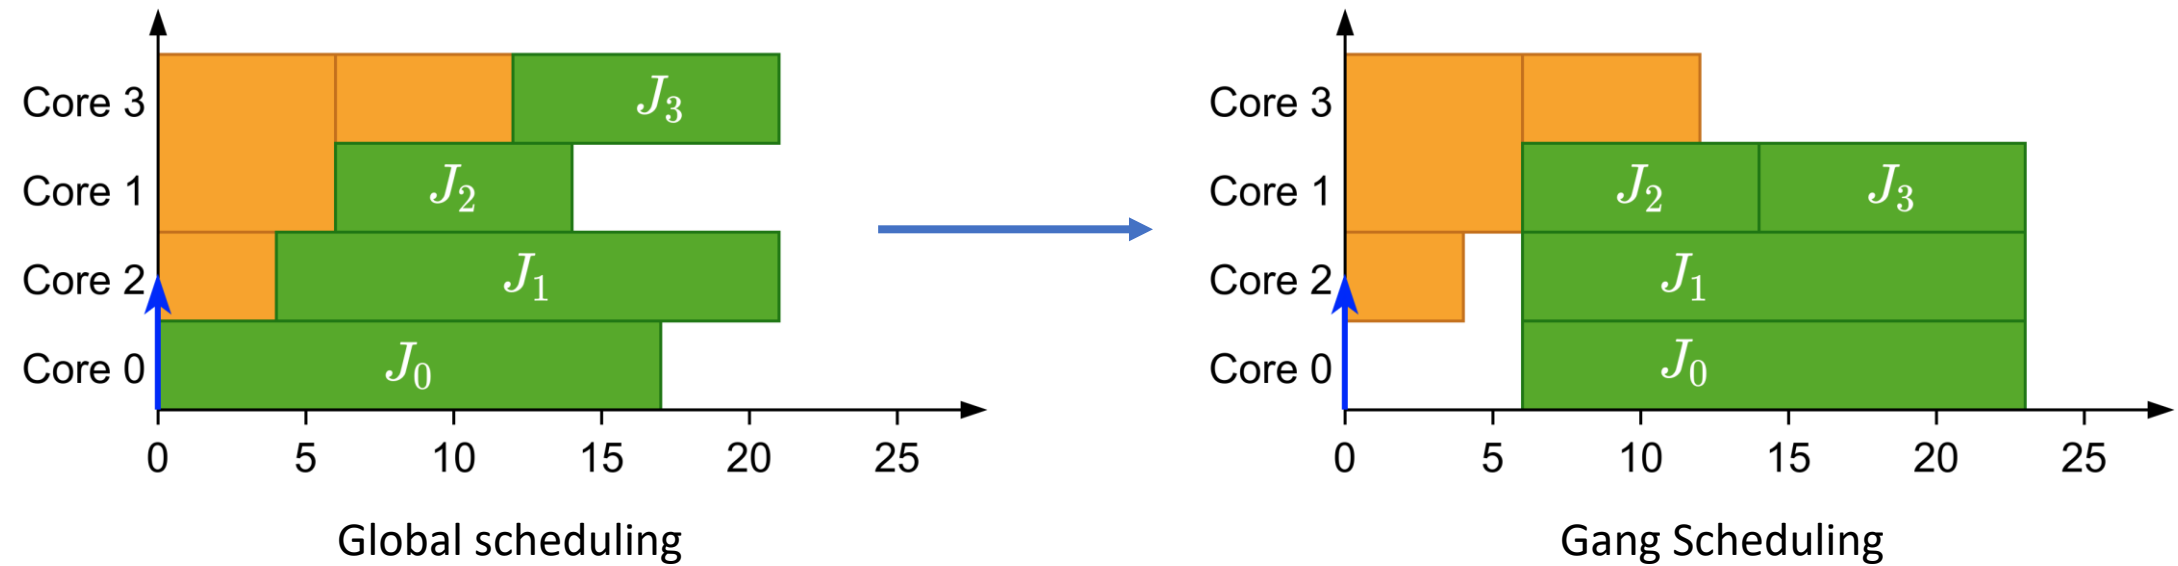
# What is gang?

# What is gang?



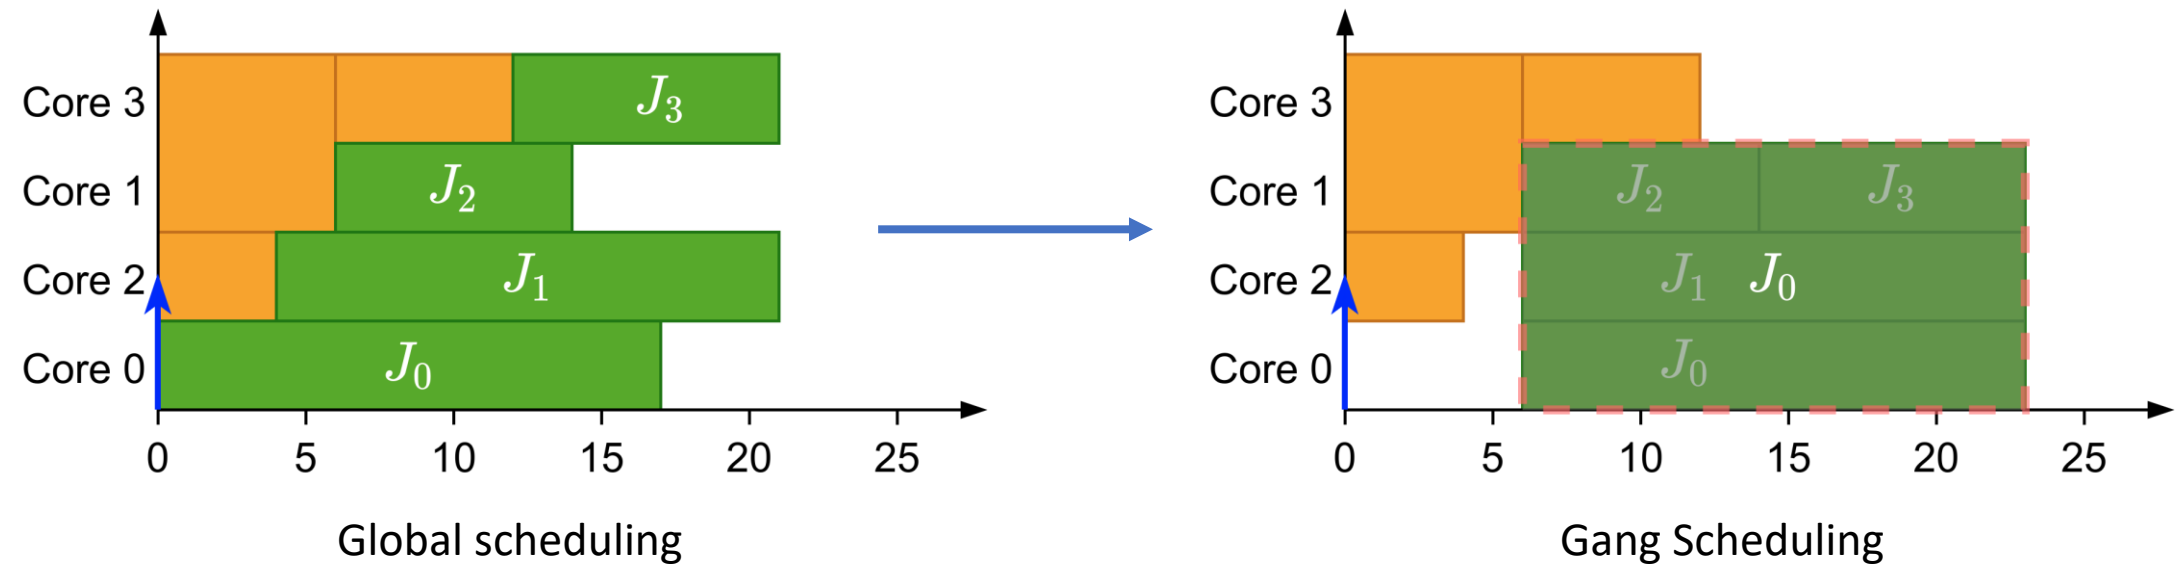
# What is gang?

- Parallel threads executed together as a “gang”




# What is gang?

- Parallel threads executed together as a “gang”
- Execution does not start until there are enough free cores




# Why gang?

# Why gang?



# Why gang?

- Efficient synchronization






# Why gang?

- Efficient synchronization
- Avoids overhead when loading initial data

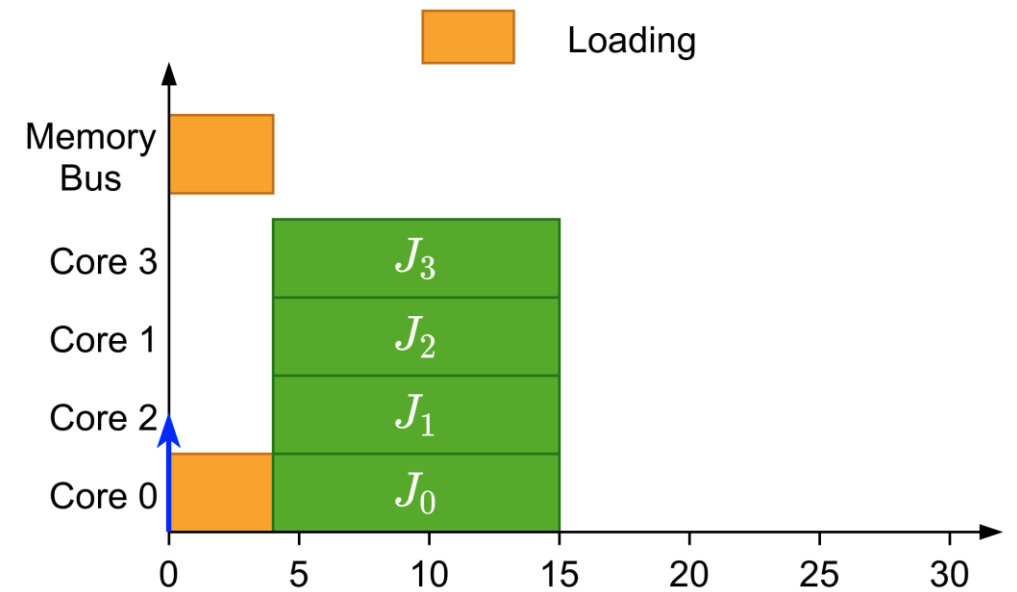

# Why gang?

- Efficient synchronization
- Avoids overhead when loading initial data



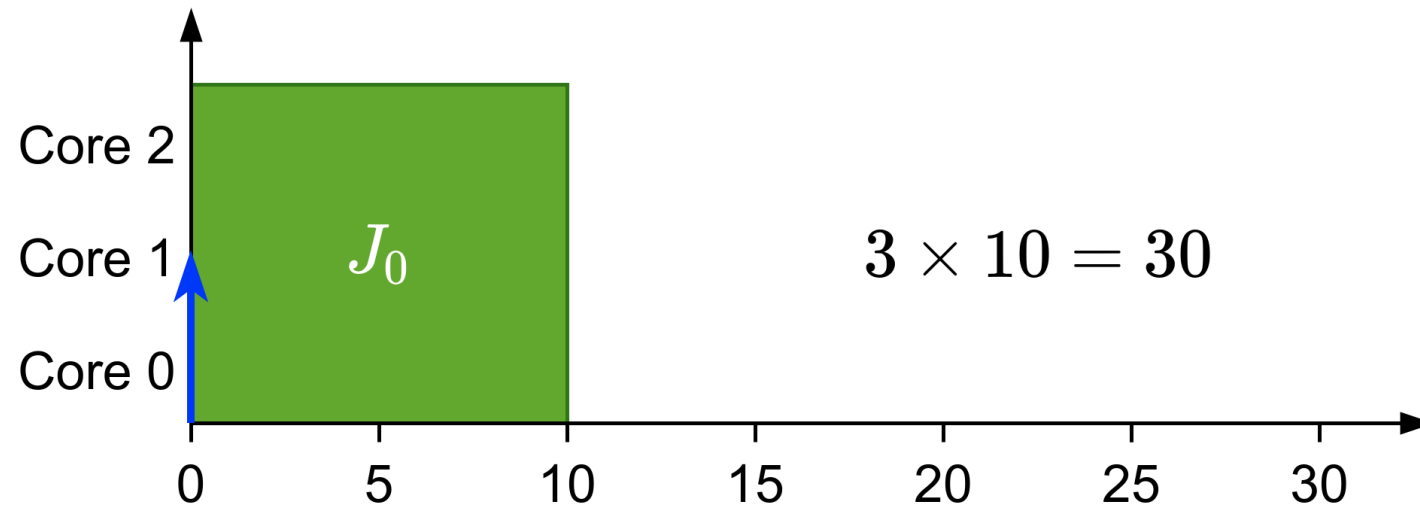
# Why gang?

- Efficient synchronization
- Avoids overhead when loading initial data



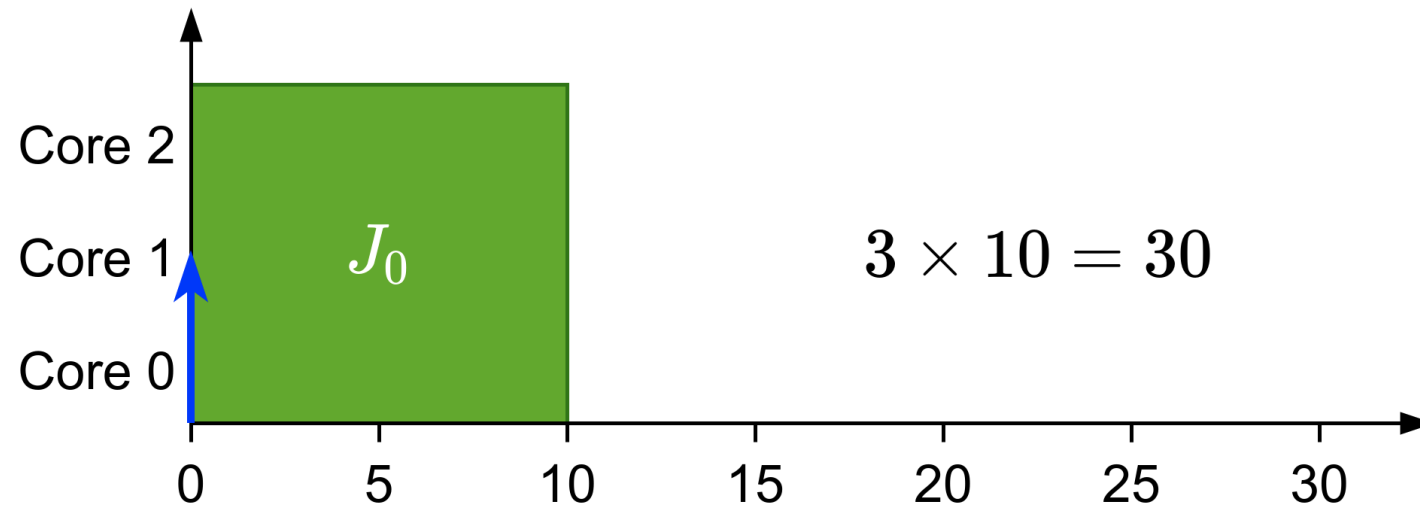
# Types of gang

# Types of gang




# Types of gang

- **Rigid:** number of cores set by programmer



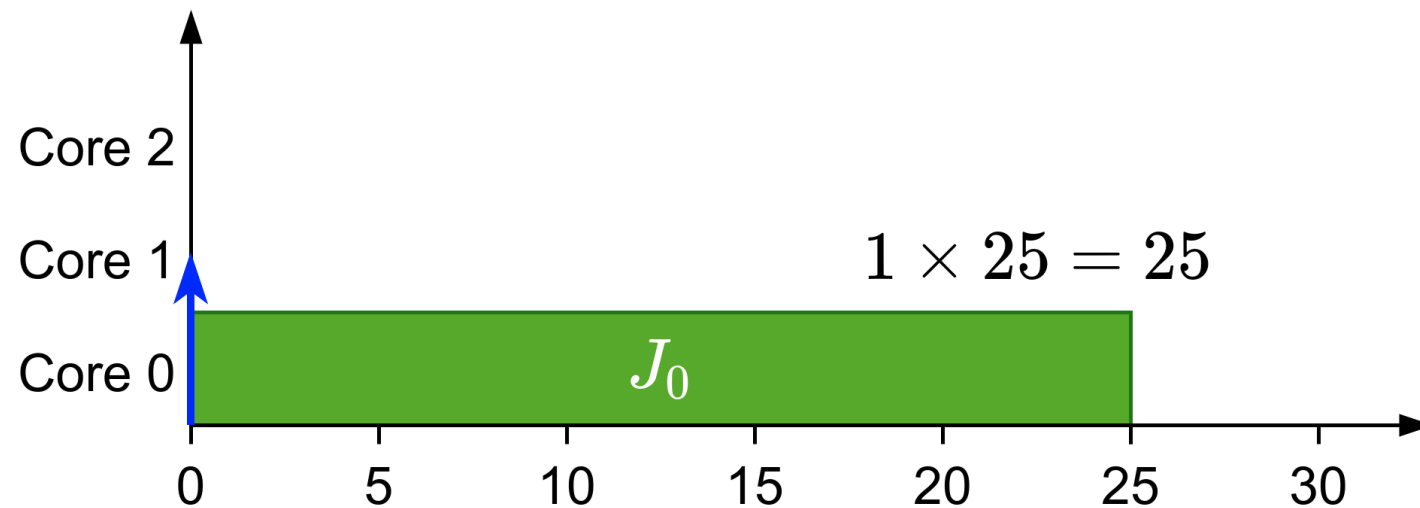
# Types of gang

- **Rigid:** number of cores set by programmer



# Types of gang

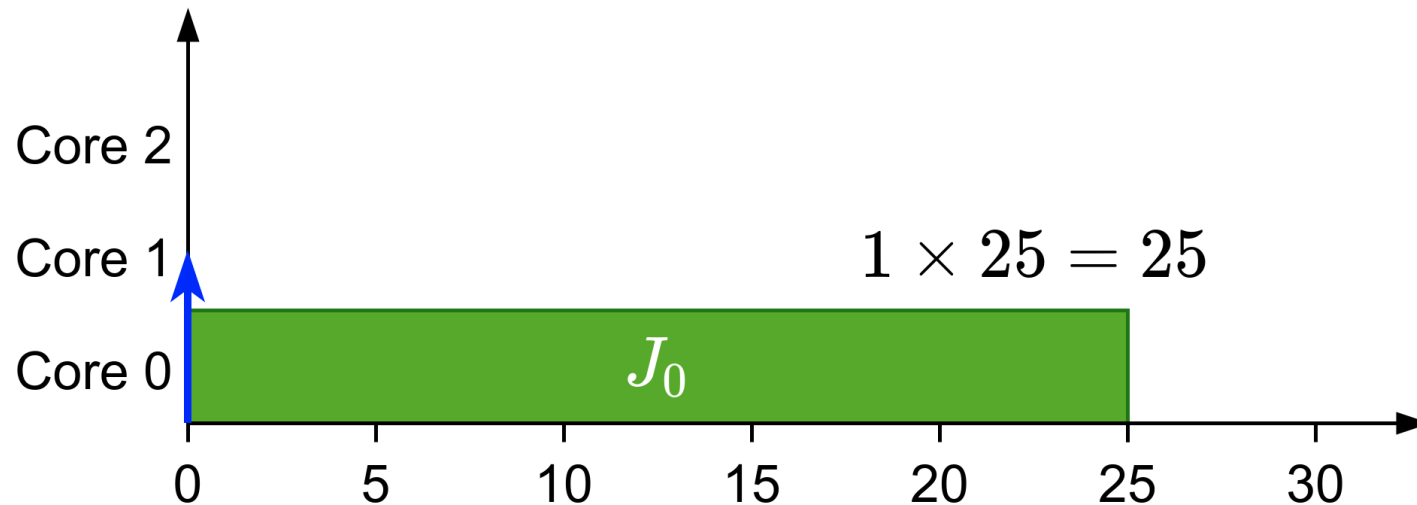
- **Rigid:** number of cores set by programmer






# Types of gang

- **Rigid**: number of cores set by programmer
- **Moldable**: number of cores assigned during scheduling




# Types of gang

- **Rigid:** number of cores set by programmer
- **Moldable:** number of cores assigned during scheduling
- **Malleable:** Number of cores can change during runtime



# Types of gang

- **Rigid**: number of cores set by programmer
- **Moldable**: number of cores assigned during scheduling
- **Malleable**: Number of cores can change during runtime



# Previous work

# Previous work

- Introduced in the context of high-performance computing<sup>[1]</sup>

# Previous work

- Introduced in the context of high-performance computing<sup>[1]</sup>
- In real-time:

# Previous work

- Introduced in the context of high-performance computing<sup>[1]</sup>
- In real-time:
  - We know that JLFP scheduler is not predictable/sustainable<sup>[2]</sup>

# Previous work


- Introduced in the context of high-performance computing<sup>[1]</sup>
- In real-time:
  - We know that JLFP scheduler is not predictable/sustainable<sup>[2]</sup>
  - Most of the work is focused in fully-preemptive solutions:
    - Optimal scheduler for rigid gang (DP-Fair)<sup>[3]</sup>
    - Moldable scheduler<sup>[4]</sup>




# Previous work

- Introduced in the context of high-performance computing<sup>[1]</sup>
- In real-time:
  - We know that JLFP scheduler is not predictable/sustainable<sup>[2]</sup>
  - Most of the work is focused in fully-preemptive solutions:
    - Optimal scheduler for rigid gang (DP-Fair)<sup>[3]</sup>
    - Moldable scheduler<sup>[4]</sup>
  - Bundled scheduling<sup>[5]</sup>
    - Tasks with precedence constraints modelled as a succession of bundles
    - Our limited-preemptive definition comes from here

# Summarizing

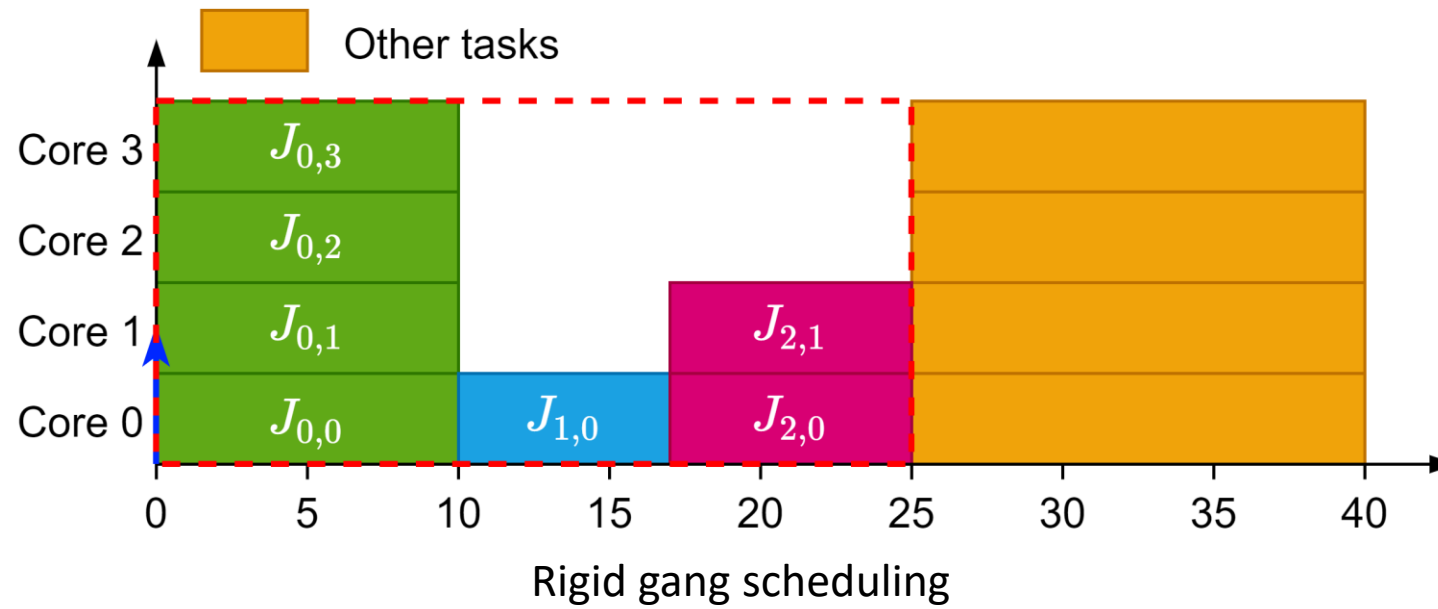


# Summarizing




# Summarizing

- Rigid gang reserves the whole block




# Summarizing

- Rigid gang reserves the whole block




# Summarizing

- Rigid gang reserves the whole block
- Bundled creates multiple rigid blocks with dependencies




# Summarizing

- Rigid gang reserves the whole block
- Bundled creates multiple rigid blocks with dependencies



# Summarizing

- Rigid gang reserves the whole block
- Bundled creates multiple rigid blocks with dependencies
- Limited-Preemptive tries to schedule these blocks in a moldable way





# Our work

# Project goals


# Project goals

- Design an accurate schedulability analysis for limited-preemptive moldable gang tasks

# Project goals


- Design an accurate schedulability analysis for limited-preemptive moldable gang tasks
- Propose a new scheduling algorithm to improve the schedulability of limited-preemptive moldable gang tasks

# Schedule Abstraction Graph




# Schedule Abstraction Graph

- Accurate and relatively fast analysis
  - Faster than an exact analysis
  - Not as pessimistic as closed-form analyses



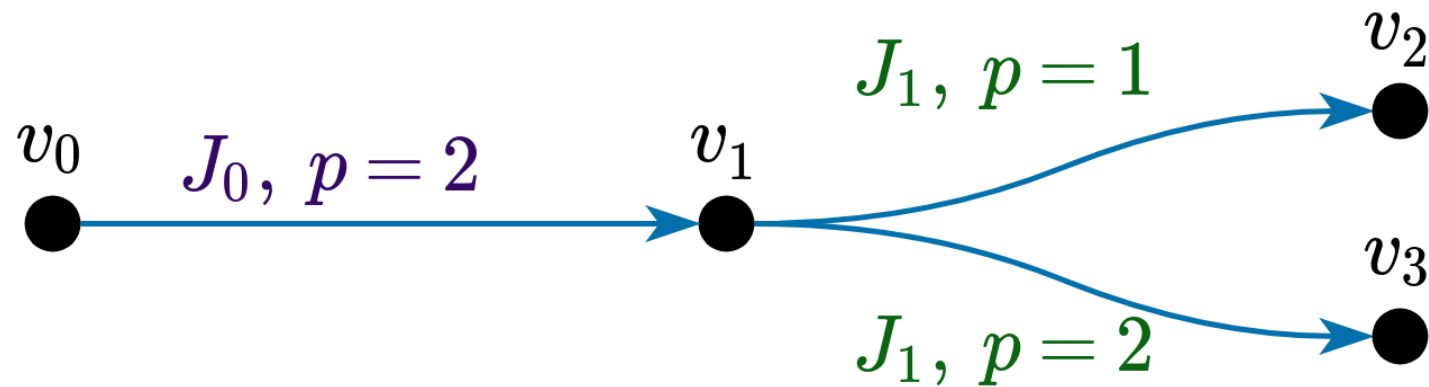
# Schedule Abstraction Graph

- Accurate and relatively fast analysis
  - Faster than an exact analysis
  - Not as pessimistic as closed-form analyses
- Models scheduler decisions



# Schedule Abstraction Graph

- Accurate and relatively fast analysis
  - Faster than an exact analysis
  - Not as pessimistic as closed-form analyses
- Models scheduler decisions
- Encodes core availability after every transition





# Job-Level Fixed Priority Scheduler for Gang

- Based on Global JLFP scheduler
- Work conserving scheduler
- Job with highest priority goes first
- Assigns maximum cores available between  $s_i^{\min}$  and  $s_i^{\max}$

# Difficulties related to SAG

- We have to consider all scenarios.
- The scheduler has to decide:
  - When to release a job
  - How many cores to assign to this job

# Analysis

# Analysis

- $A_c^{\min}$  time at which we have  $c$  cores possibly available
- $A_c^{\max}$  time at which we have  $c$  cores certainly available

# Analysis

- $A_c^{\min}$  time at which we have  $c$  cores possibly available
- $A_c^{\max}$  time at which we have  $c$  cores certainly available
- $EST_i$  Earliest Start Time
- $LST_i$  Latest Start Time
- $EFT_i$  Earliest Finishing Time
- $LFT_i$  Latest Finishing Time

# Analysis

- $A_c^{\min}$  time at which we have  $c$  cores possibly available
- $A_c^{\max}$  time at which we have  $c$  cores certainly available
- $EST_i$  Earliest Start Time
- $LST_i$  Latest Start Time
- $EFT_i$  Earliest Finishing Time
- $LFT_i$  Latest Finishing Time

$$EST_i \leq LST_i$$

# Analysis

- $A_c^{\min}$  time at which we have  $c$  cores possibly available
- $A_c^{\max}$  time at which we have  $c$  cores certainly available
- $EST_i^p$  Earliest Start Time
- $LST_i^p$  Latest Start Time
- $EFT_i^p$  Earliest Finishing Time
- $LFT_i^p$  Latest Finishing Time

$$EST_i^p \leq LST_i^p$$

# Analysis

$$EST_i^p = \max\{r_i^{\min}, A_p^{\min}\}$$

- Job cannot start before
  - Being released
  - Enough cores are available

$$LST_i^p = \min\{t_{p+1}, t_{wc}, t_{high} - 1\}$$

- Job cannot start with  $p$  cores after:
  - $p + 1$  cores are available as JLFP would schedule it with  $p + 1$  cores
  - A lower priority task is ready because JLFP is work-conserving
  - A higher priority task is ready



# Analysis

- Obtain  $EFT_i^p$  and  $LFT_i^p$  from:

$$EFT_i^p = EST_i^p + c_i^{\min}(p)$$


$$LFT_i^p = LST_i^p + c_i^{\max}(p)$$

- And compute new  $A_c^{\min}$  and  $A_c^{\max}$

# New scheduler


# LPMRGS

# LPMRGS



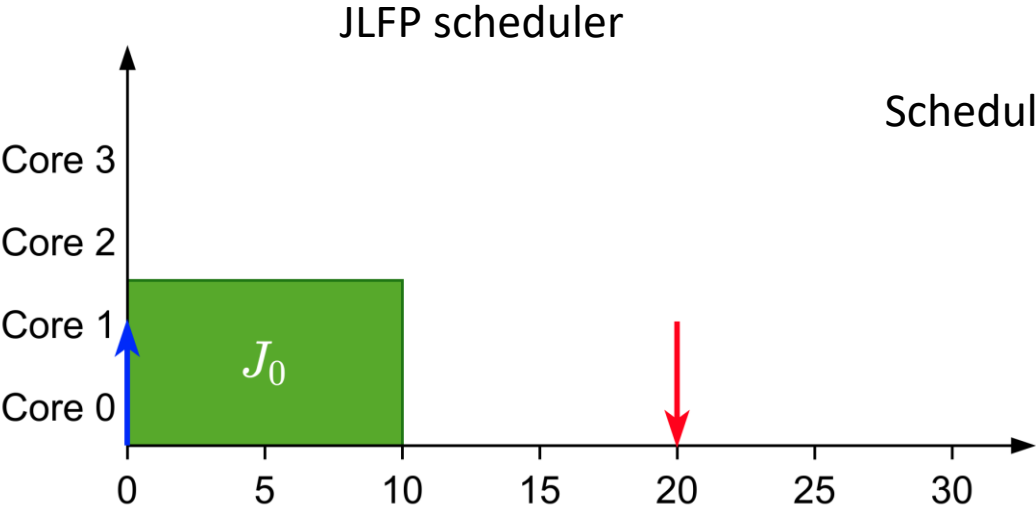
	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS




JLFP scheduler

Scheduling:  $J_1$


	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20


# LPMRGS



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20




# LPMRGS



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS




	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS

- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks

# LPMRGS


- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS


- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS


- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS


- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS

- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks




	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20



# LPMRGS


- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS


- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# LPMRGS

- Limited-Preemptive Moldable Reservation Gang Scheduler
- Non-work conserving scheduler
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	$s_j^{\min}$	$s_j^{\max}$	$d_i$	$c_i(v)$
$J_0$	2	2	100	10
$J_1$	3	3	20	5
$J_2$	1	1	100	20
$J_3$	1	3	100	24, 13, 20

# Questions?