

Schedulability analysis of limited-preemptive moldable gang tasks

Joan Marcè i Igual



Daily Supervisor

Co-supervisor

Geoffrey Nelissen

Mitra Nasri

27th of August, 2020

Real-time systems

Real-time systems

- Systems of which correctness does depends not only on **logical** results but also on **timing constraints**

Real-time systems

- Systems of which correctness does depends not only on **logical** results but also on **timing constraints**



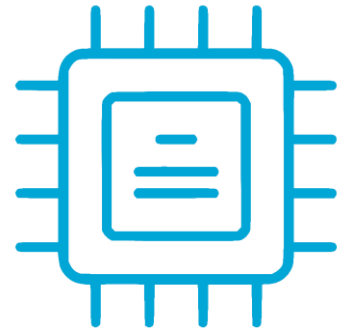
Real-time systems

- Systems of which correctness does depends not only on **logical** results but also on **timing constraints**



Real-time systems

- Systems of which correctness does depends not only on **logical** results but also on **timing constraints**
- Multicore systems



Definitions

Definitions

- Task
 - A functionality of the system

Definitions

- Task
 - A functionality of the system
- Job
 - Instance of a task

Definitions

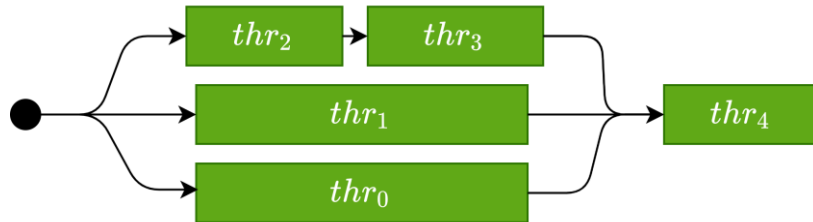
- Task
 - A functionality of the system
- Job
 - Instance of a task
- Schedule
 - A particular assignment of jobs to the processors and time intervals

Definitions


- Task
 - A functionality of the system
- Job
 - Instance of a task
- Schedule
 - A particular assignment of jobs to the processors and time intervals
- Scheduling policy
 - Algorithm that produces a schedule
 - FIFO, Round-Robin, JLFP, EDF

What is gang?


What is gang?



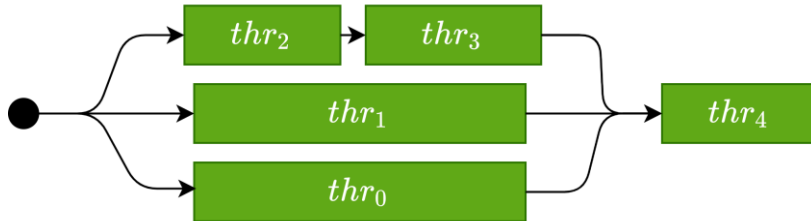
What is gang?



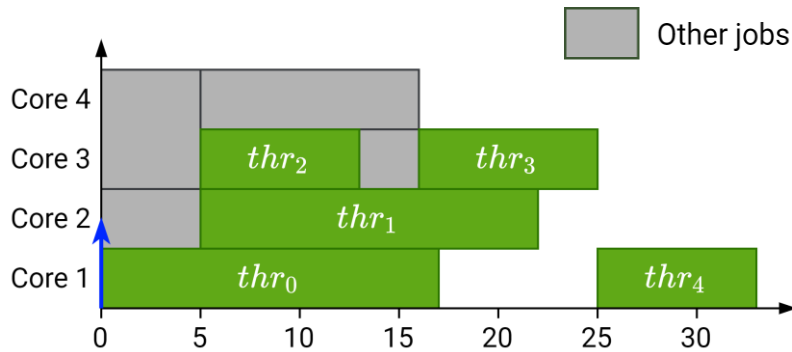
Global scheduling



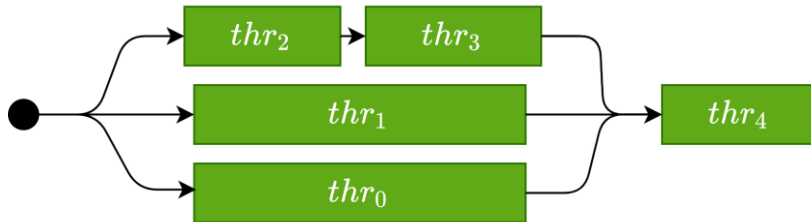
What is gang?



Global scheduling




What is gang?




Parallel threads together as a “gang”

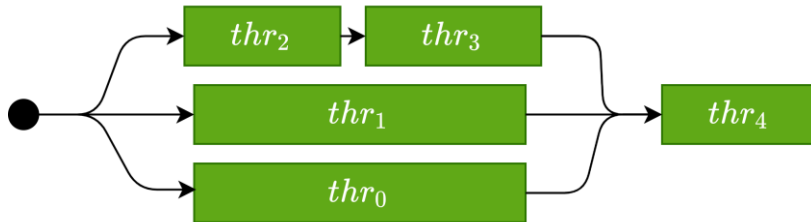
Global scheduling



Gang Scheduling

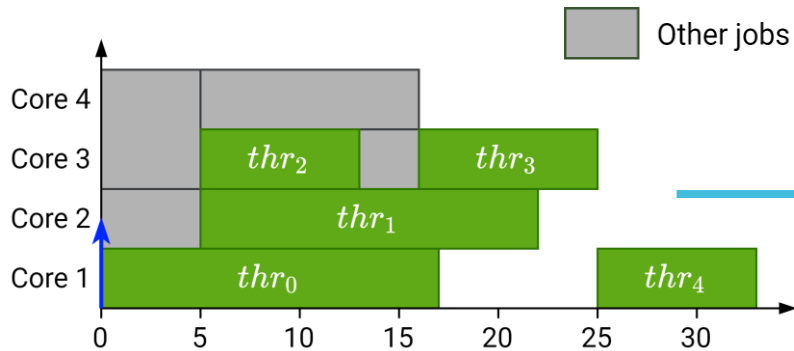


What is gang?

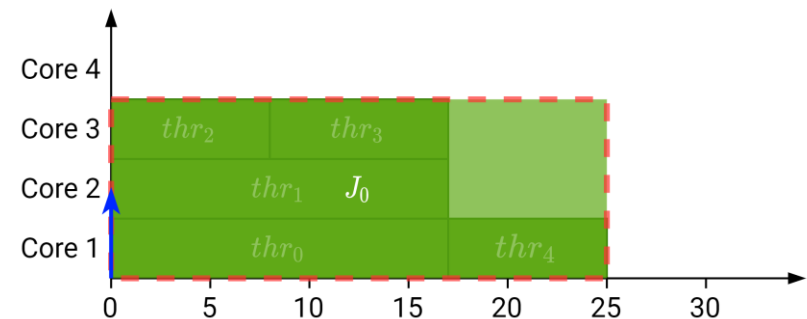


Parallel threads together as a “gang”

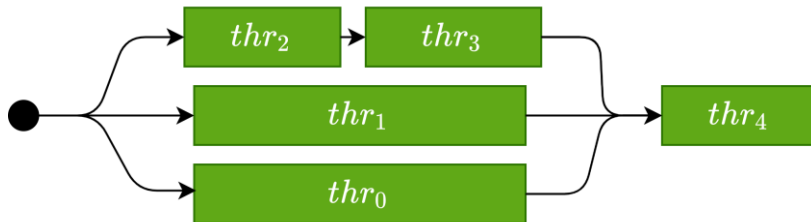
Global scheduling



Gang Scheduling



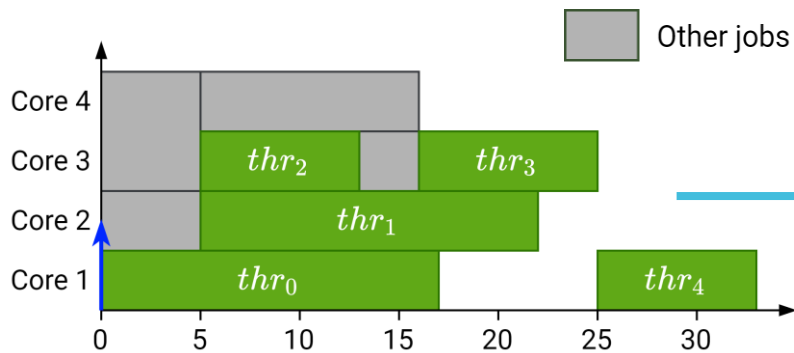
What is gang?



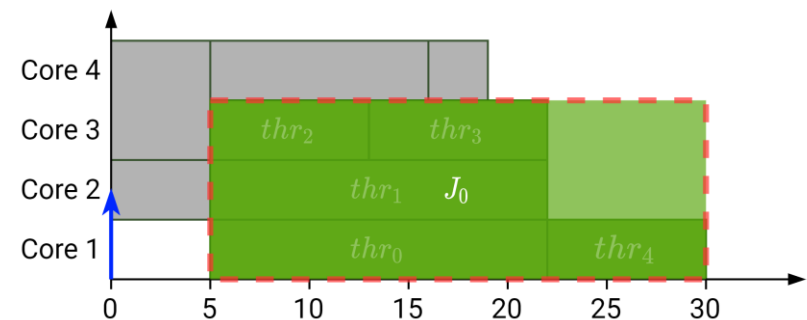
Parallel threads together as a “gang”

Execution does not start until there are enough cores

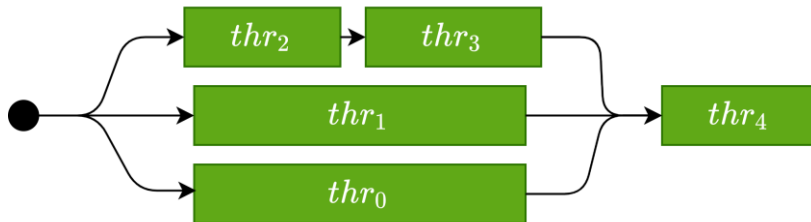
Global scheduling



Gang Scheduling



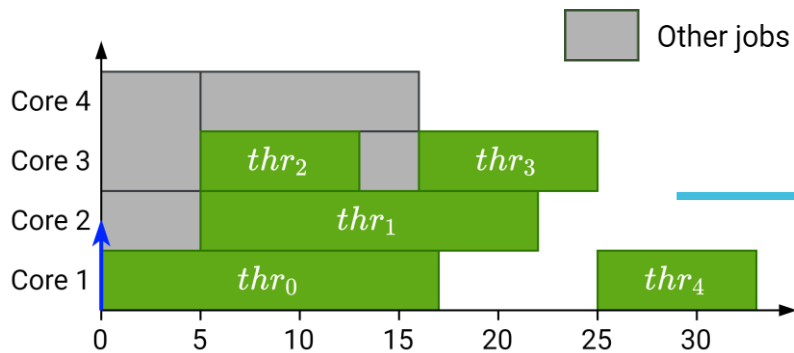
What is gang?




Parallel threads together as a “gang”

Execution does not start until there are enough cores


Global scheduling



Gang Scheduling




What is gang?




Parallel threads together as a “gang”

Execution does not start until there are enough cores

Global scheduling



Gang Scheduling



Why gang?

Why gang?

- More efficient synchronization

Why gang?


- More efficient synchronization
- Reduces variability in the execution

Why gang?

- More efficient synchronization
- Reduces variability in the execution
- Avoids overhead when loading initial data

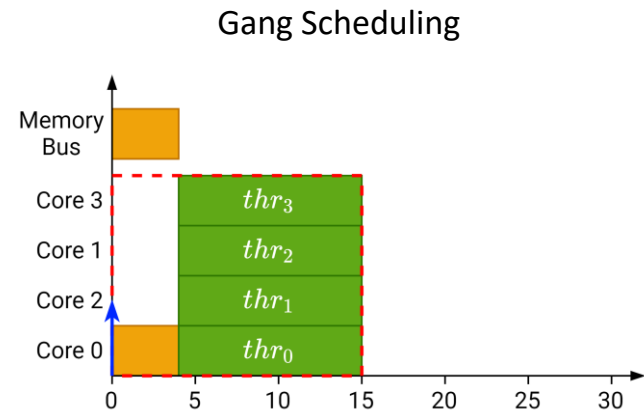
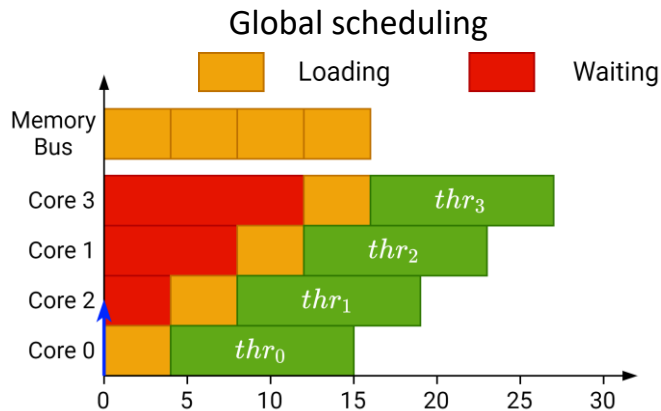
Why gang?

- More efficient synchronization
- Reduces variability in the execution
- Avoids overhead when loading initial data



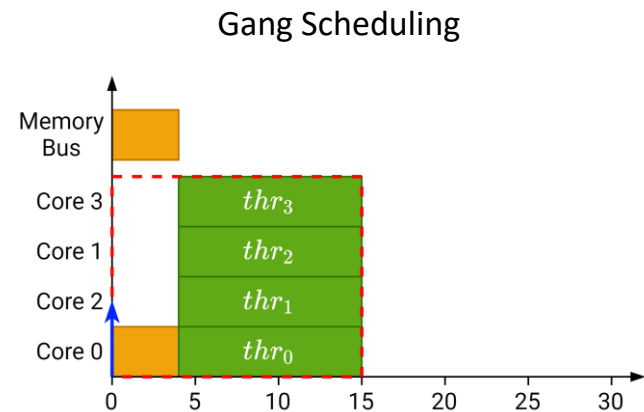
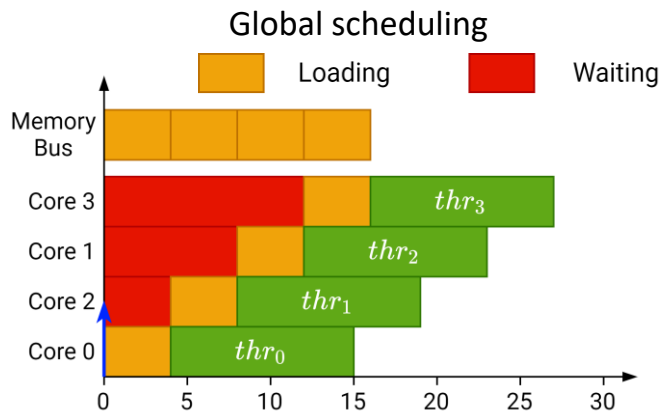
Why gang?

- More efficient synchronization
- Reduces variability in the execution
- Avoids overhead when loading initial data



Why gang?

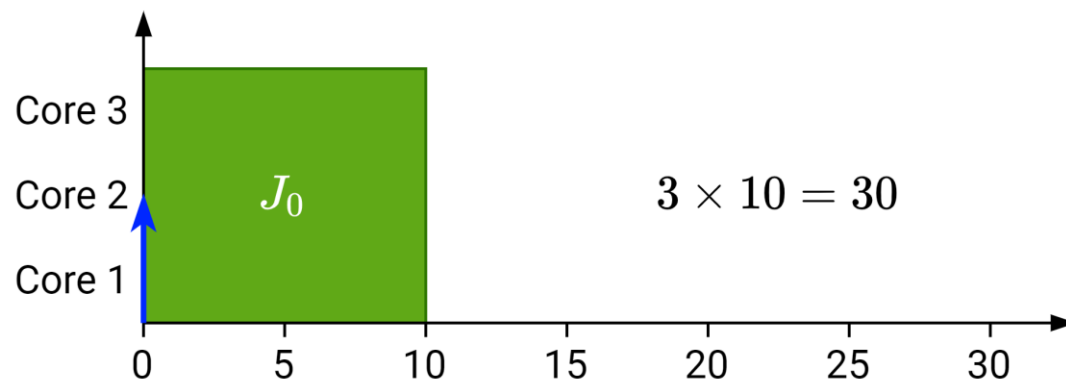
- More efficient synchronization
- Reduces variability in the execution
- Avoids overhead when loading initial data
- Shows its full potential when executed non-preemptively



Types of gang

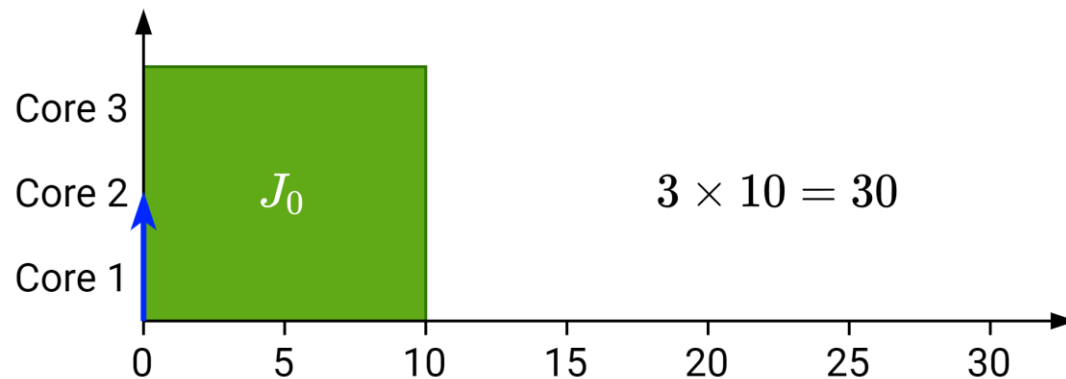
Types of gang

- **Rigid:** number of cores set by programmer




Types of gang

- **Rigid**: number of cores set by programmer
- **Moldable**: number of cores assigned when job is dispatched



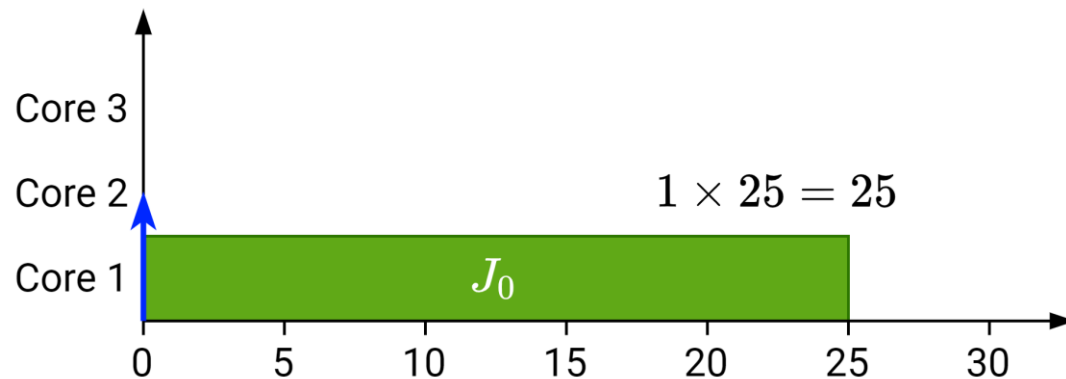
Types of gang

- **Rigid:** number of cores set by programmer
- **Moldable:** number of cores assigned when job is dispatched



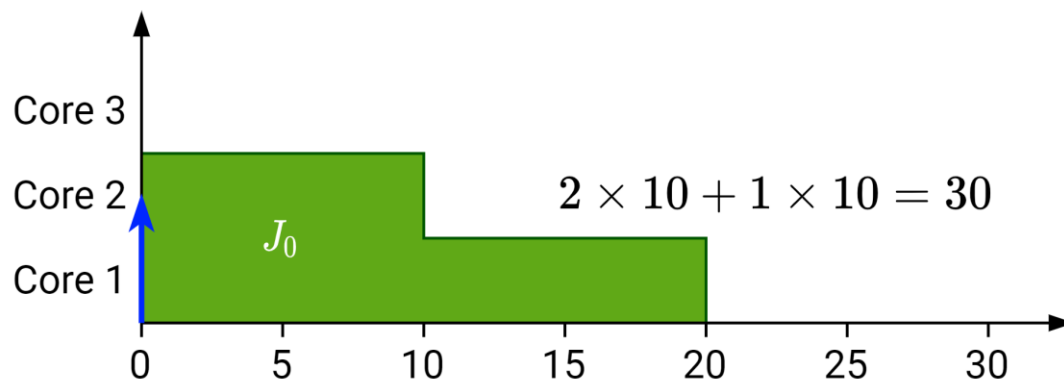
Types of gang

- **Rigid**: number of cores set by programmer
- **Moldable**: number of cores assigned when job is dispatched




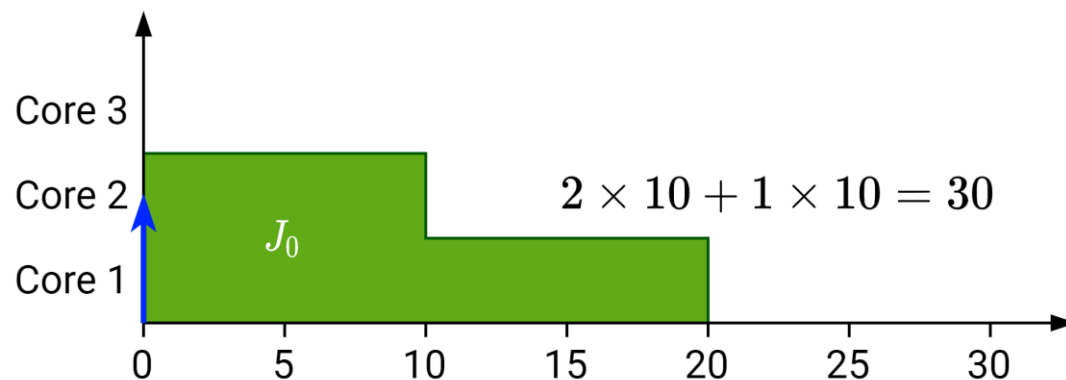
Types of gang

- **Rigid:** number of cores set by programmer
- **Moldable:** number of cores assigned when job is dispatched
- **Malleable:** number of cores can change during runtime



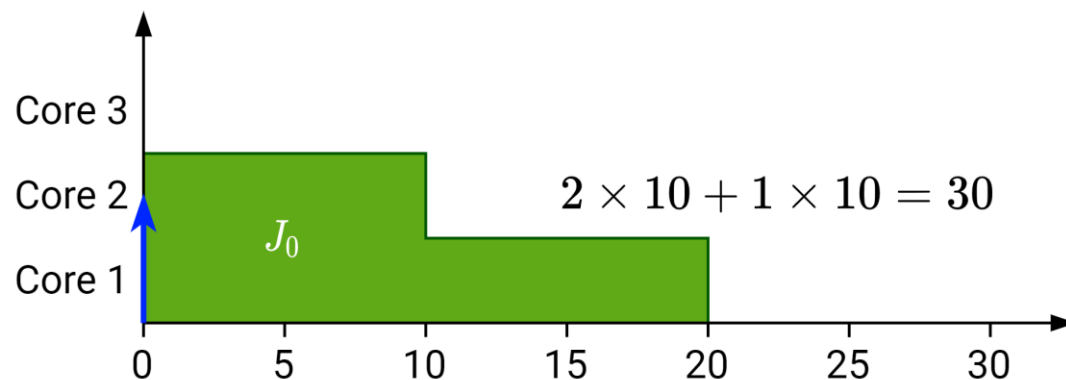
Types of gang

- **Rigid**: number of cores set by programmer
- **Moldable**: number of cores assigned when job is dispatched
- **Malleable**: number of cores can change during runtime →  Hard to implement



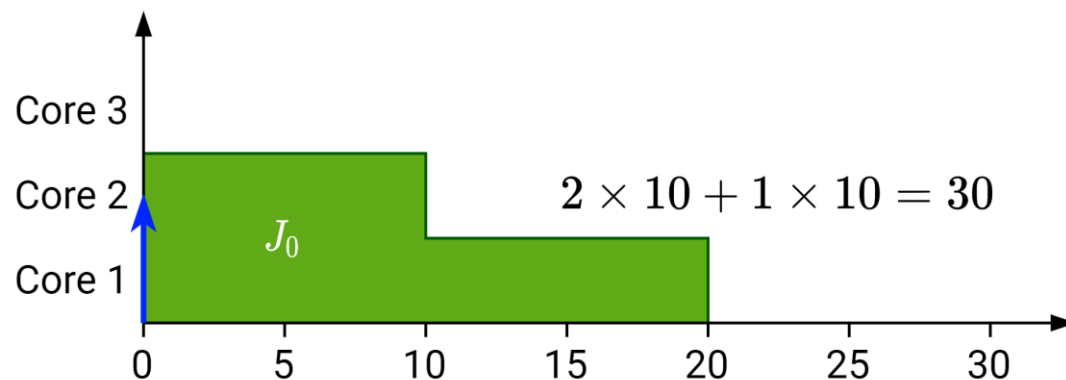
Types of gang

- **Rigid**: number of cores set by programmer → 🗨️ Wastes resources
- **Moldable**: number of cores assigned when job is dispatched
- **Malleable**: number of cores can change during runtime → 🗨️ Hard to implement



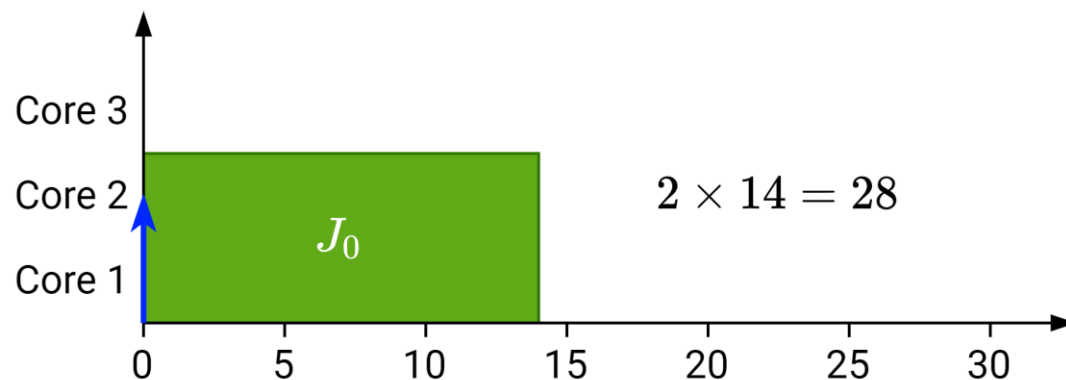
Types of gang

- **Rigid**: number of cores set by programmer → 🗨️ Wastes resources
- **Moldable**: number of cores assigned when job is dispatched → 👍 Flexibility
- **Malleable**: number of cores can change during runtime → 🗨️ Hard to implement

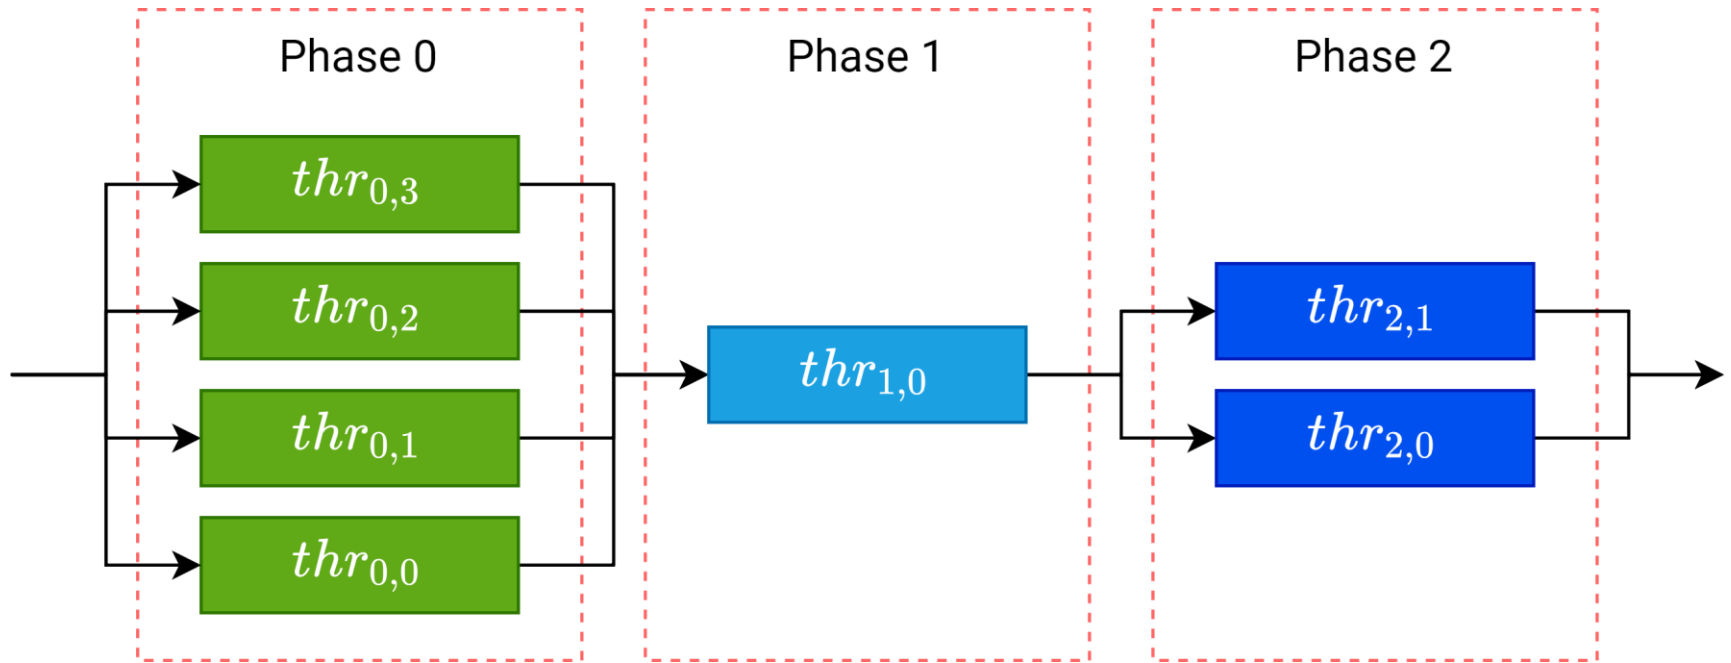


Types of gang

- **Rigid**: number of cores set by programmer → 🗑️ Wastes resources
- **Moldable**: number of cores assigned when job is dispatched → 👍 Flexibility
- **Malleable**: number of cores can change during runtime → 🗑️ Hard to implement





Bundled scheduling^[1] vs limited-preemptive



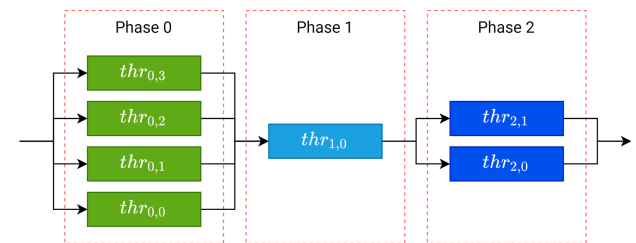
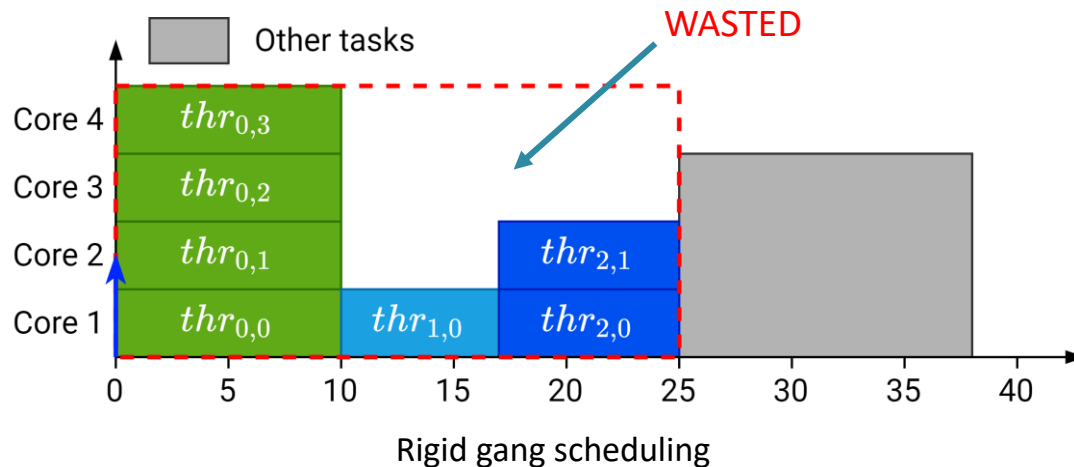
Bundled scheduling^[1] vs limited-preemptive

- Rigid gang reserves the whole block




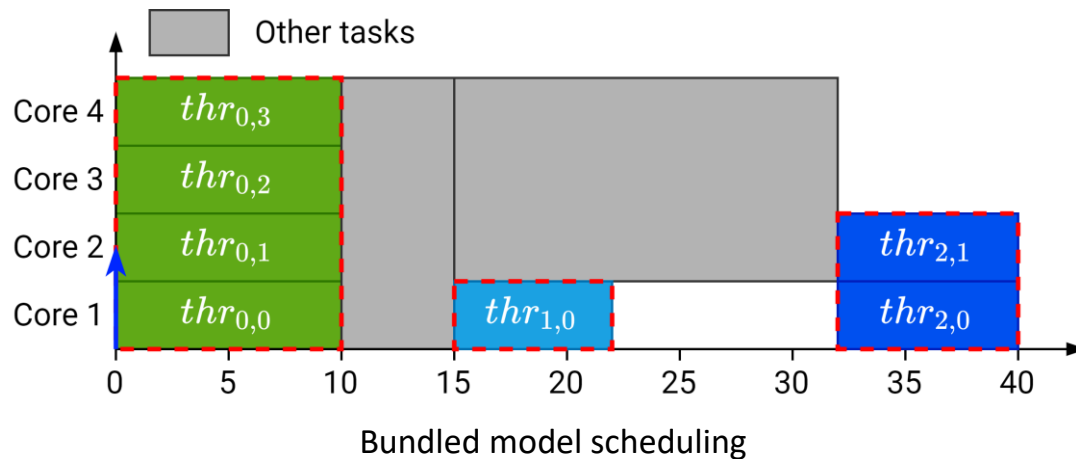
Bundled scheduling^[1] vs limited-preemptive

- Rigid gang reserves the whole block



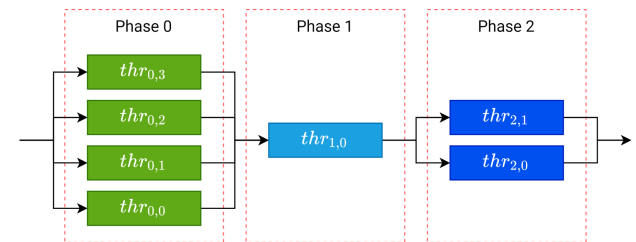
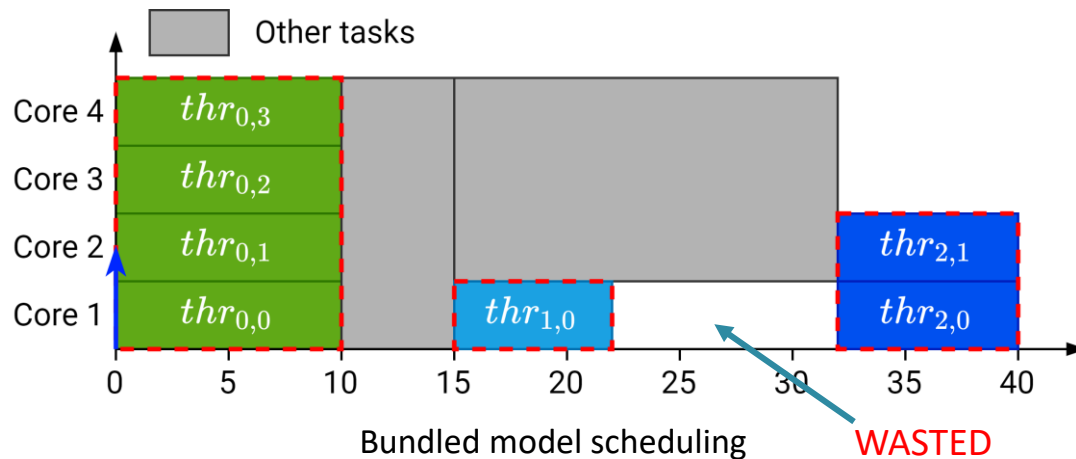
Bundled scheduling^[1] vs limited-preemptive

- Rigid gang reserves the whole block
- Bundled creates **rigid blocks** with dependencies



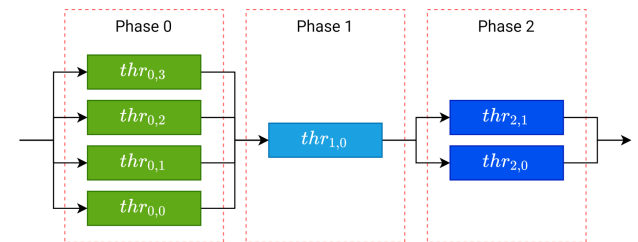
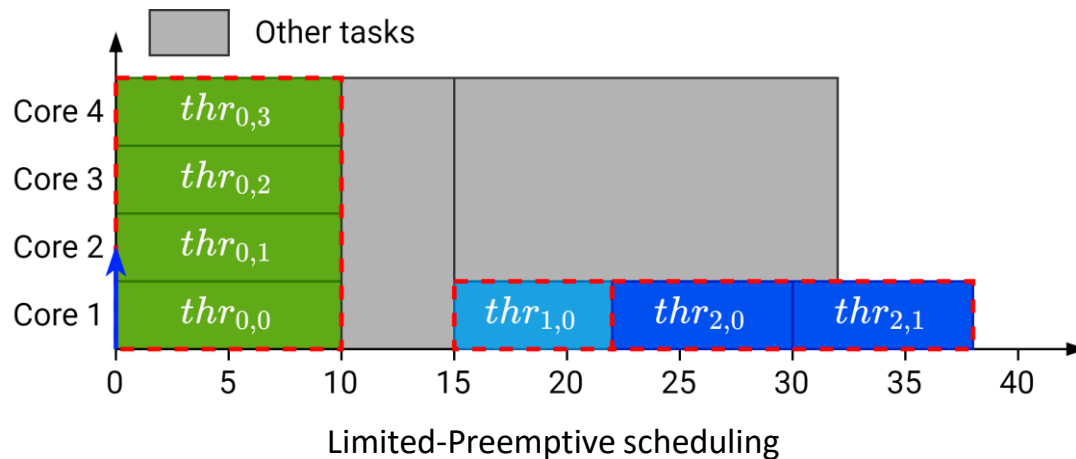
Bundled scheduling^[1] vs limited-preemptive

- Rigid gang reserves the whole block
- Bundled creates **rigid blocks** with dependencies



Bundled scheduling^[1] vs limited-preemptive

- Rigid gang reserves the whole block
- Bundled creates **rigid blocks** with dependencies
- Limited-Preemptive creates **moldable blocks** with dependencies



Job-level fixed-priority scheduling (JLFP) for gang

Job-level fixed-priority scheduling (JLFP) for gang

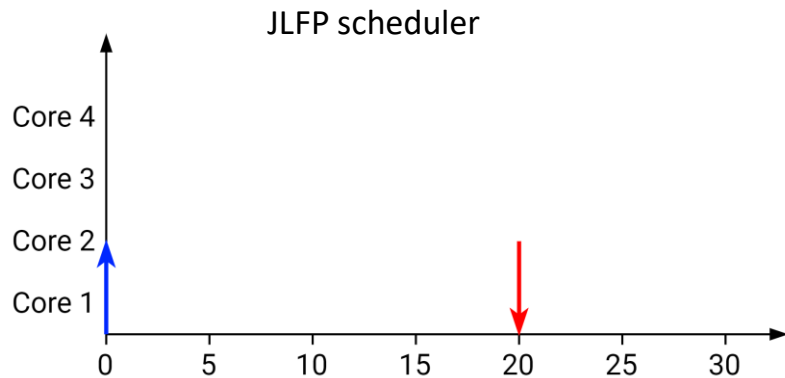
- Based on global JLFP scheduler

Job-level fixed-priority scheduling (JLFP) for gang

- Based on global JLFP scheduler
- Assigns maximum number of available cores to a job between the range of cores that the job allows

Job-level fixed-priority scheduling (JLFP) for gang

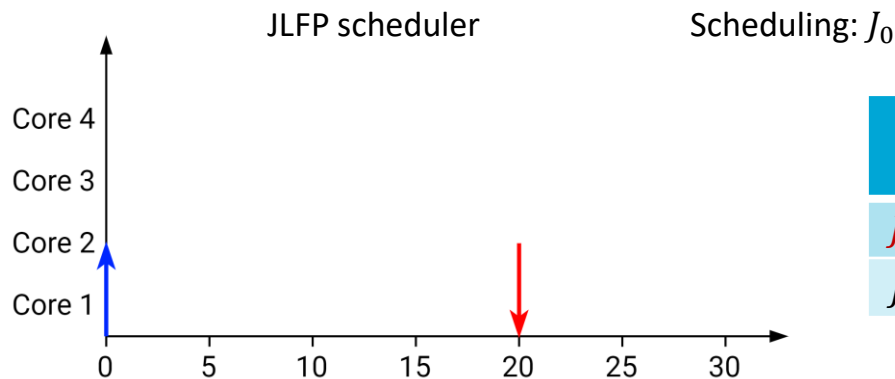
- Based on global JLFP scheduler
- Assigns maximum number of available cores to a job between the range of cores that the job allows



	Priority	Min cores	Max cores	Deadline	Execution time
J_0	High	2	3	∞	15, 10
J_1	Low	2	2	20	15

Job-level fixed-priority scheduling (JLFP) for gang

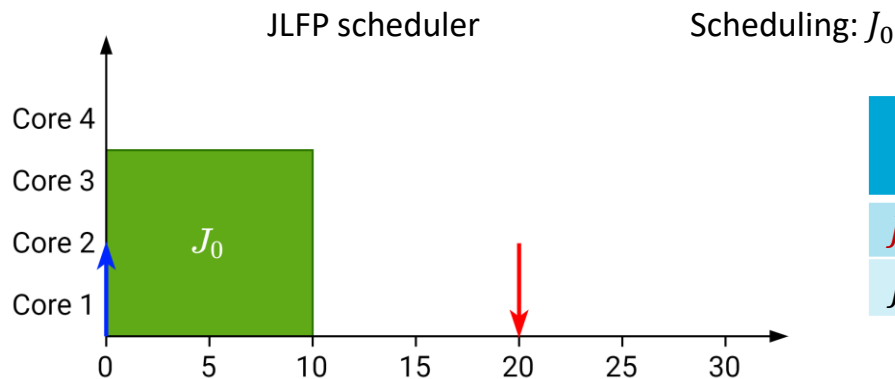
- Based on global JLFP scheduler
- Assigns maximum number of available cores to a job between the range of cores that the job allows



	Priority	Min cores	Max cores	Deadline	Execution time
J_0	High	2	3	∞	15, 10
J_1	Low	2	2	20	15

Job-level fixed-priority scheduling (JLFP) for gang

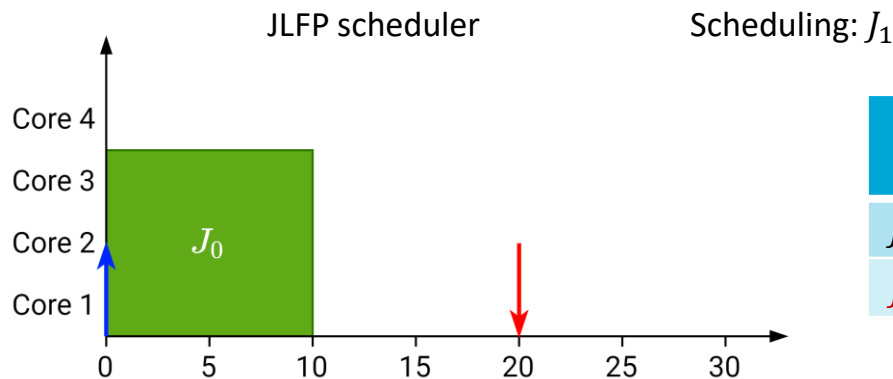
- Based on global JLFP scheduler
- Assigns maximum number of available cores to a job between the range of cores that the job allows



	Priority	Min cores	Max cores	Deadline	Execution time
J_0	High	2	3	∞	15, 10
J_1	Low	2	2	20	15

Job-level fixed-priority scheduling (JLFP) for gang

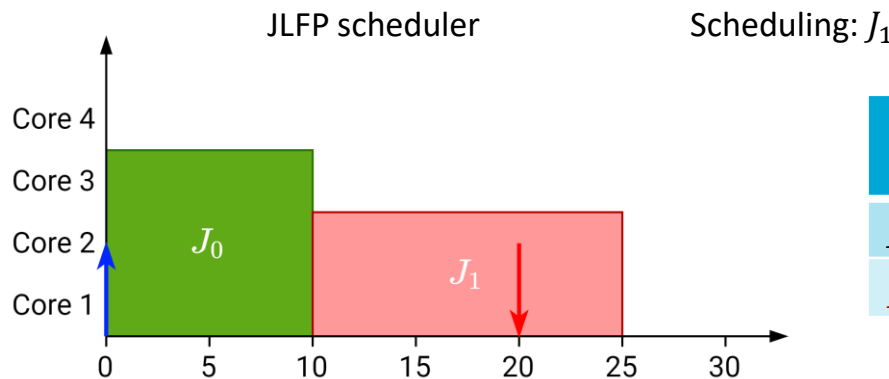
- Based on global JLFP scheduler
- Assigns maximum number of available cores to a job between the range of cores that the job allows



	Priority	Min cores	Max cores	Deadline	Execution time
J_0	High	2	3	∞	15, 10
J_1	Low	2	2	20	15

Job-level fixed-priority scheduling (JLFP) for gang

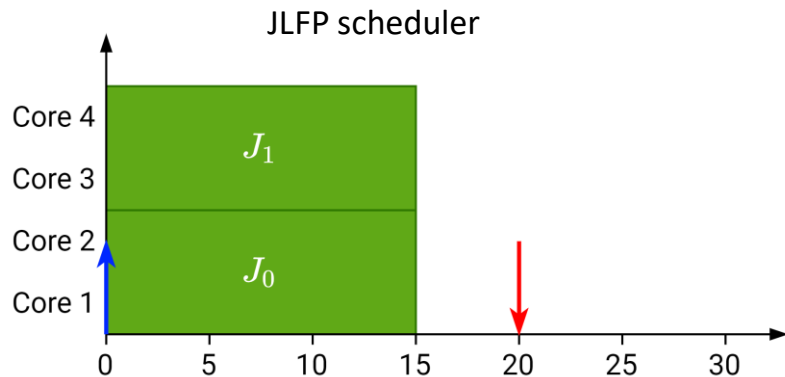
- Based on global JLFP scheduler
- Assigns maximum number of available cores to a job between the range of cores that the job allows



	Priority	Min cores	Max cores	Deadline	Execution time
J_0	High	2	3	∞	15, 10
J_1	Low	2	2	20	15

Job-level fixed-priority scheduling (JLFP) for gang

- Based on global JLFP scheduler
- Assigns maximum number of available cores to a job between the range of cores that the job allows



	Priority	Min cores	Max cores	Deadline	Execution time
J_0	High	2	3	∞	15, 10
J_1	Low	2	2	20	15

Previous work

Previous work

Introduced in high-performance
computing in 1982^[1]

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

^[1]Ousterhout, 1982

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

^[1]Ousterhout, 1982

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

- Job-level fixed-priority^[2]

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

- Job-level fixed-priority^[2]
- Earliest deadline first^[3]

^[1]Ousterhout, 1982

^[3]Kato et al., 2009

^[2]Goossens et al., 2010

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

- Job-level fixed-priority^[2]
- Earliest deadline first^[3]

Schedulers

^[1]Ousterhout, 1982

^[2]Goossens et al., 2010

^[3]Kato et al., 2009

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

- Job-level fixed-priority^[2]
- Earliest deadline first^[3]

Schedulers

- Optimal for rigid gang (DP-Fair)^[4]

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

- Job-level fixed-priority^[2]
- Earliest deadline first^[3]

Schedulers

- Optimal for rigid gang (DP-Fair)^[4]
- Moldable gang^[5]

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

- Job-level fixed-priority^[2]
- Earliest deadline first^[3]

Schedulers

- Optimal for rigid gang (DP-Fair)^[4]
- Moldable gang^[5]

Non-preemptive solutions

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

- Job-level fixed-priority^[2]
- Earliest deadline first^[3]

Schedulers

- Optimal for rigid gang (DP-Fair)^[4]
- Moldable gang^[5]

Non-preemptive solutions

Schedulability tests

Previous work

Introduced in high-performance computing in 1982^[1]

Preemptive solutions

Schedulability tests

- Job-level fixed-priority^[2]
- Earliest deadline first^[3]

Schedulers

- Optimal for rigid gang (DP-Fair)^[4]
- Moldable gang^[5]

Non-preemptive solutions

Schedulability tests

- Earliest deadline first for rigid gang^[6]

Our work

Project goals

Project goals

1. Design an accurate schedulability **analysis** for limited-preemptive **moldable gang tasks**

Project goals

1. Design an accurate schedulability **analysis** for limited-preemptive **moldable gang tasks**
2. Evaluate the impact of the level of parallelism assigned to the jobs

Project goals

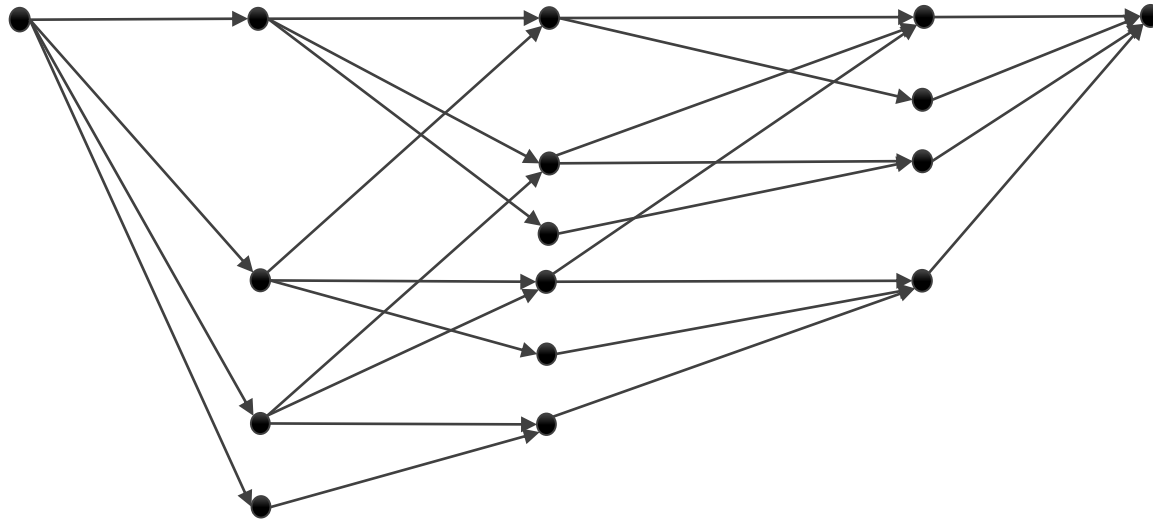
1. Design an accurate schedulability **analysis** for limited-preemptive **moldable gang tasks**
2. Evaluate the impact of the level of parallelism assigned to the jobs
3. Propose a **new scheduling algorithm** to improve the schedulability of limited-preemptive **moldable gang tasks**

Agenda


- Gang schedulability analysis
- New scheduling policy

Schedule abstraction graph


Schedule abstraction graph



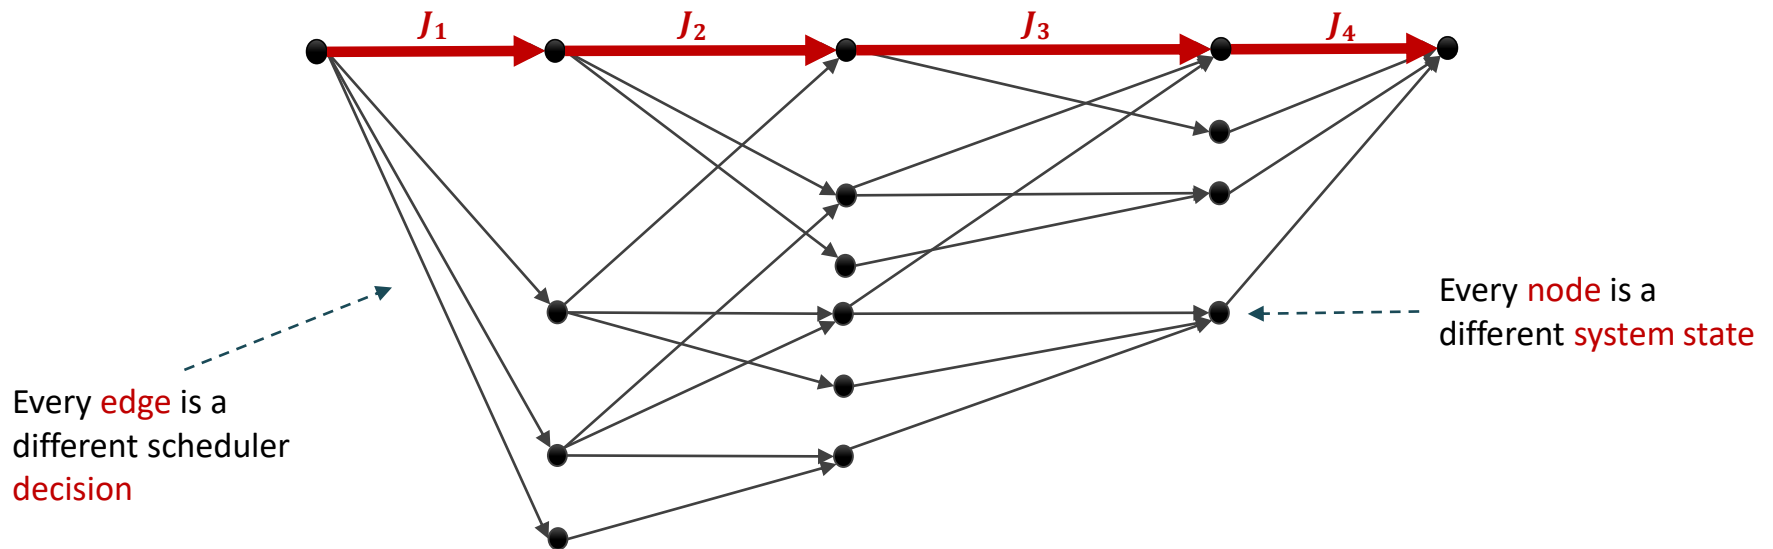
Schedule abstraction graph



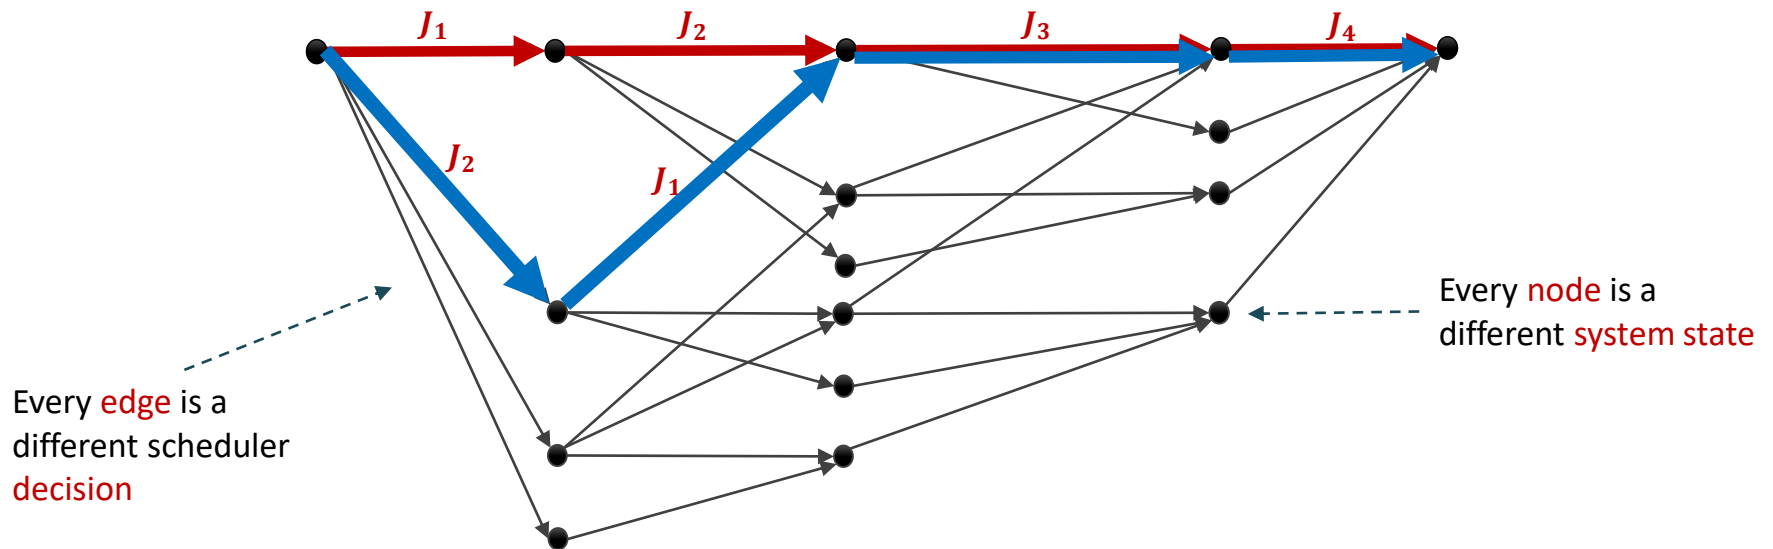
Schedule abstraction graph



Schedule abstraction graph

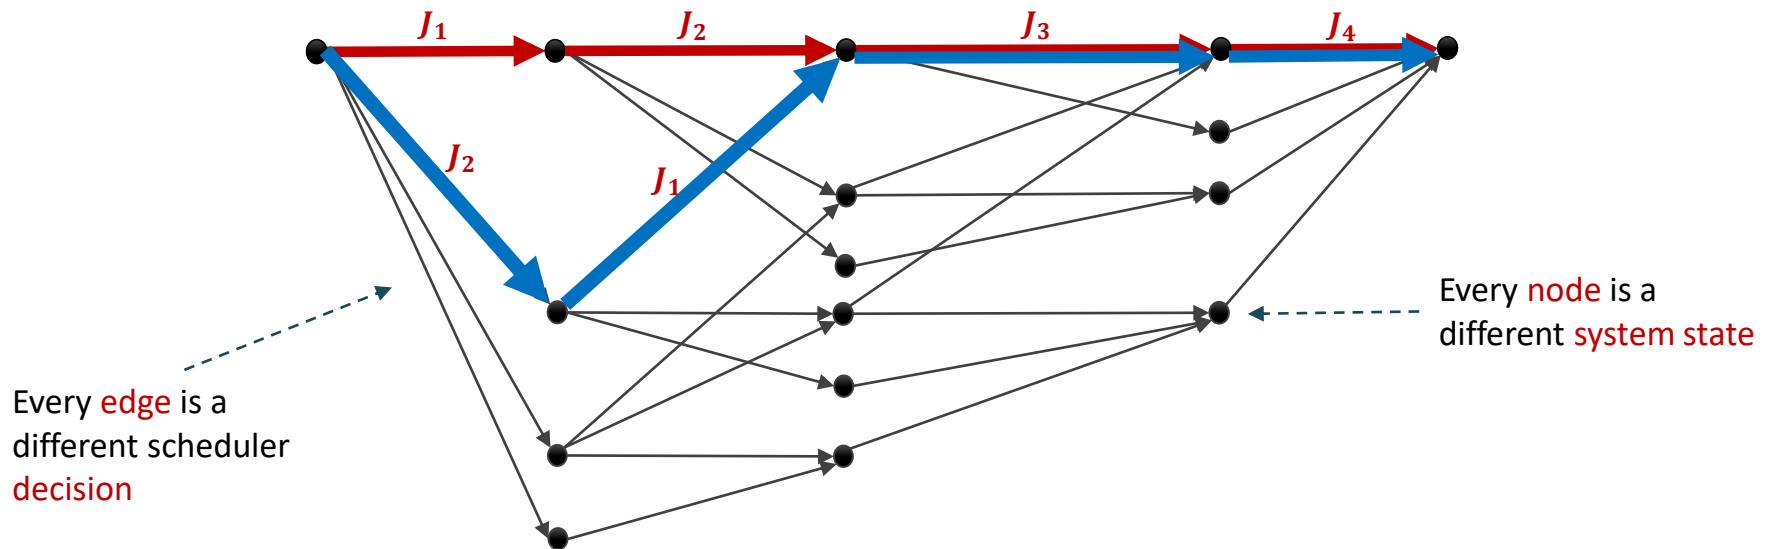


Schedule abstraction graph



Schedule abstraction graph

- It is a technique that allows:
 - Search for all possible schedules
 - Aggregate “similar” schedules



Extending SAG

Extending SAG

- Update system state representation

Extending SAG

- Update system state representation
- Update expansion rules

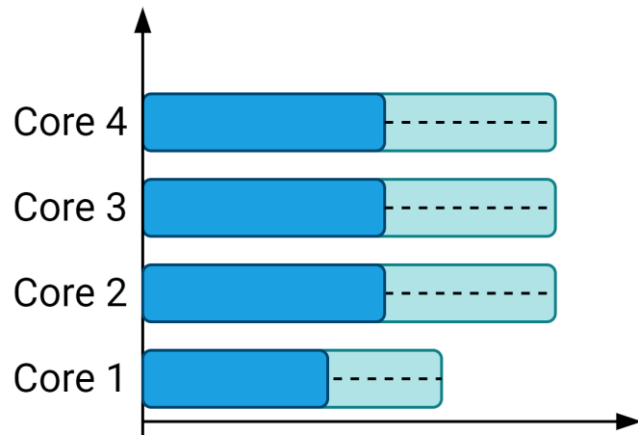
Extending SAG

- Update system state representation
- Update expansion rules
- Update merge rules

New state representation

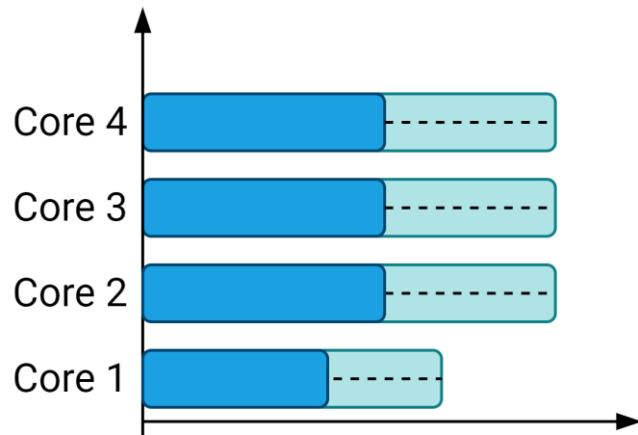
New state representation

- We already had core availabilities



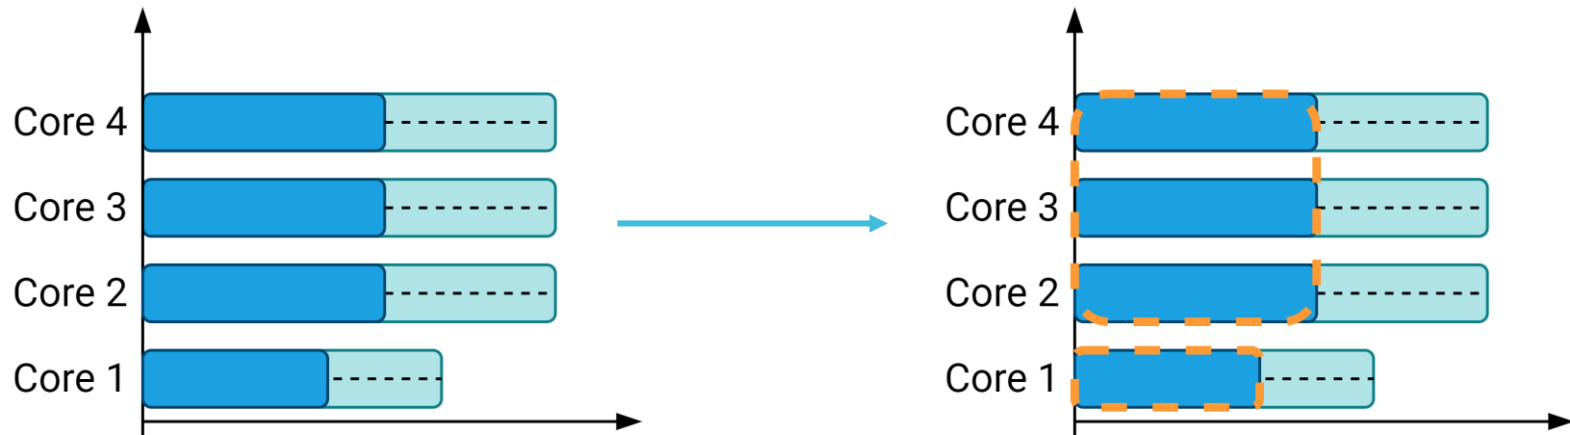
New state representation

- We already had core availabilities
- We need to keep track of



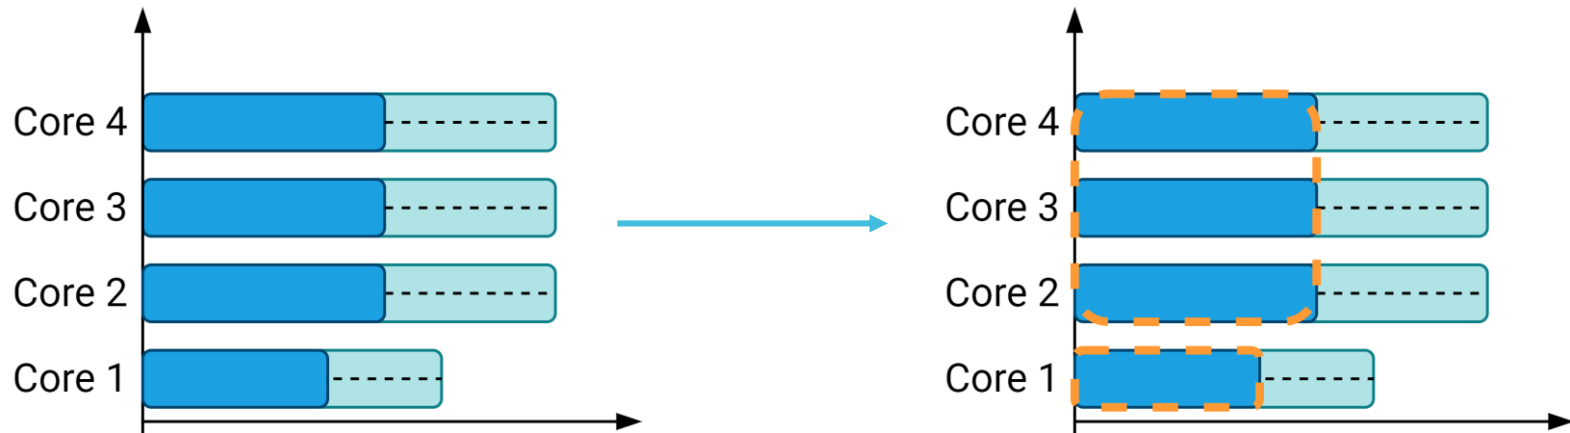
New state representation

- We already had core availabilities
- We need to keep track of
 - Groups of cores



New state representation

- We already had core availabilities
- We need to keep track of
 - Groups of cores
 - Certainly running jobs



New SAG expansion rules

New SAG expansion rules

- Previously a state was created for every schedulable job


New SAG expansion rules

- Previously a state was created for every schedulable job

s_0

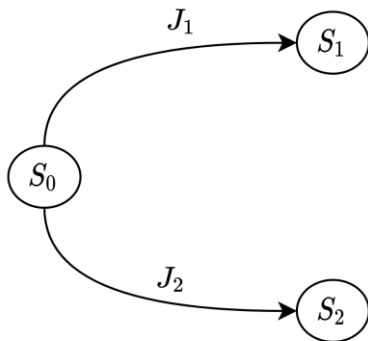
New SAG expansion rules

- Previously a state was created for every schedulable job




New SAG expansion rules

- Previously a state was created for every schedulable job




New SAG expansion rules

- Previously a state was created for every schedulable job
- Now a state is created for every job and possible number of cores

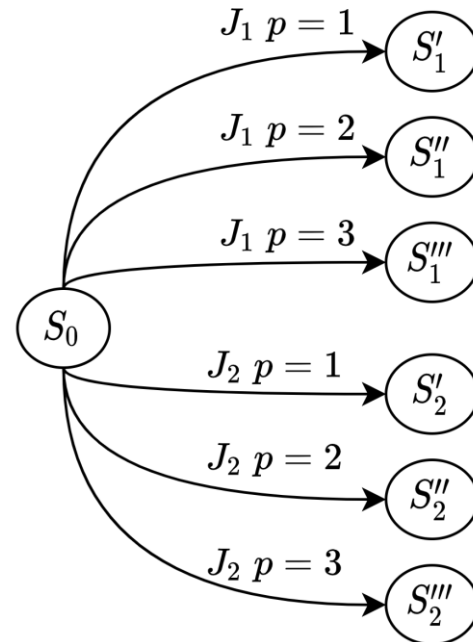


New SAG expansion rules

- Previously a state was created for every schedulable job
- Now a state is created for every job and possible number of cores
- Can stimulate state-space explosion

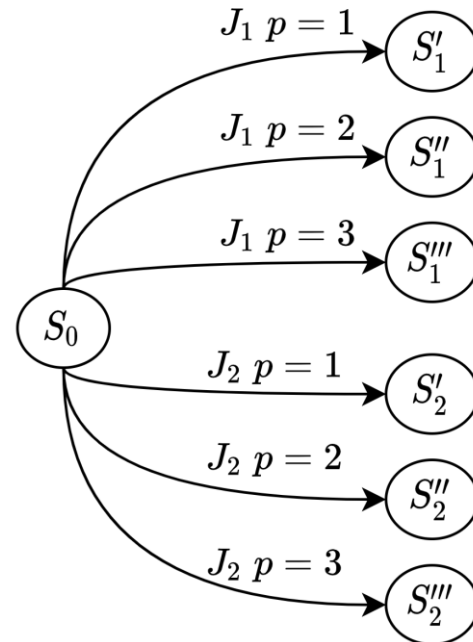


New SAG expansion rules




New SAG expansion rules

- Purge states by checking that



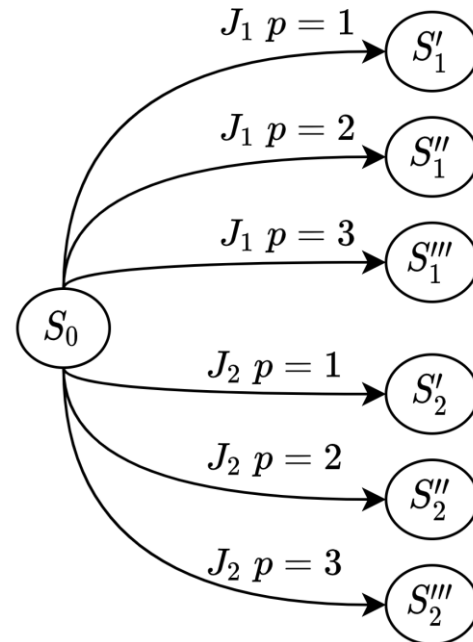
New SAG expansion rules

- Purge states by checking that
 - p cores available
 -




New SAG expansion rules

- Purge states by checking that
 - p cores available
 - More cores **not** available



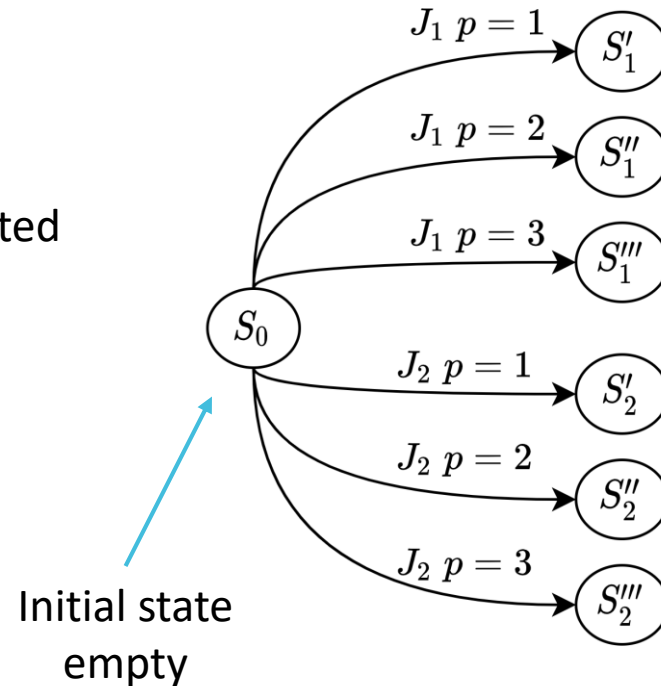
New SAG expansion rules

- Purge states by checking that
 - p cores available
 - More cores **not** available
 - Precedence constraints are respected



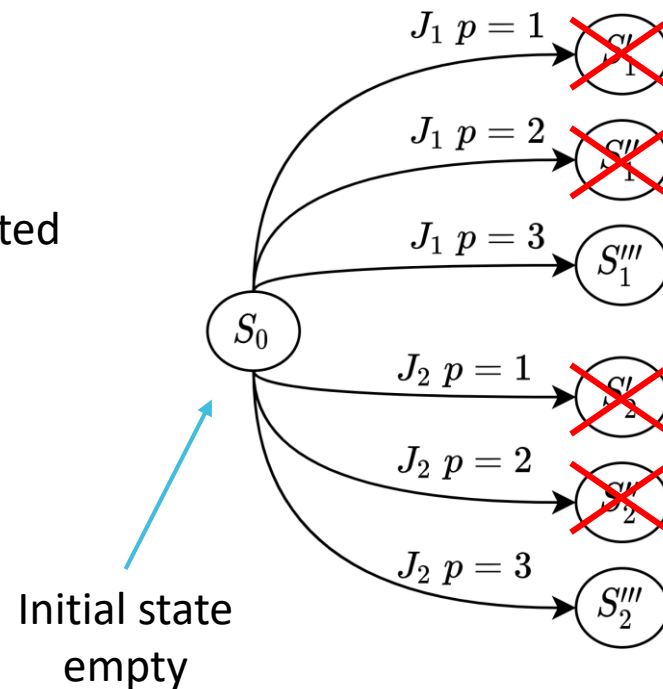
New SAG expansion rules

- Purge states by checking that
 - p cores available
 - More cores **not** available
 - Precedence constraints are respected



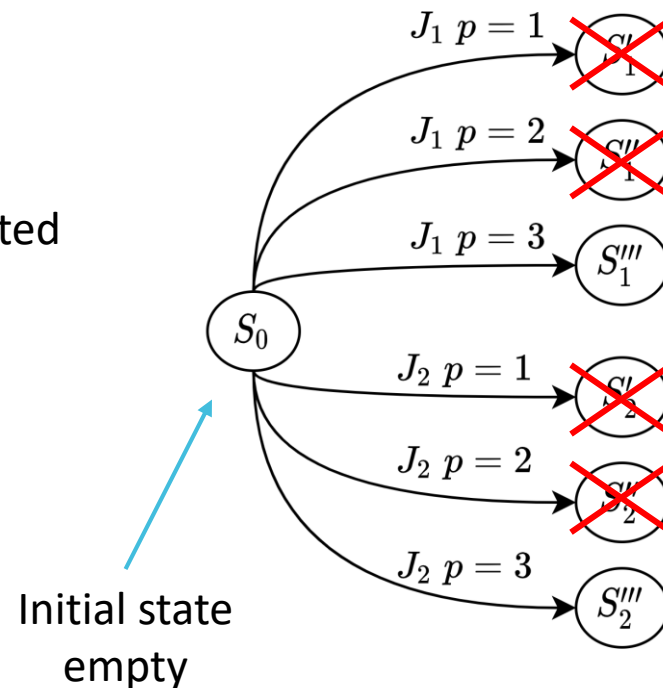
New SAG expansion rules

- Purge states by checking that
 - p cores available
 - More cores **not** available
 - Precedence constraints are respected



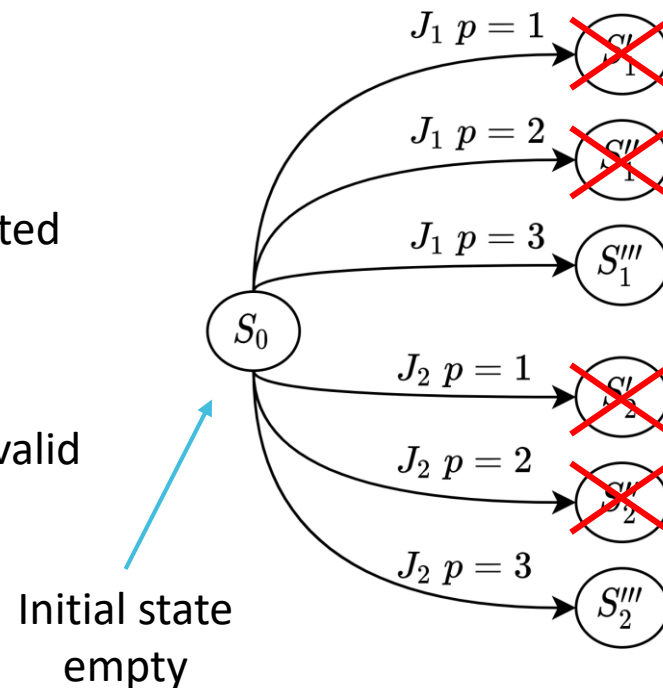
New SAG expansion rules

- Purge states by checking that
 - p cores available
 - More cores **not** available
 - Precedence constraints are respected
- Exploring more states



New SAG expansion rules

- Purge states by checking that
 - p cores available
 - More cores **not** available
 - Precedence constraints are respected
- Exploring more states
 - 👍 Is **safe**, does not make analysis invalid

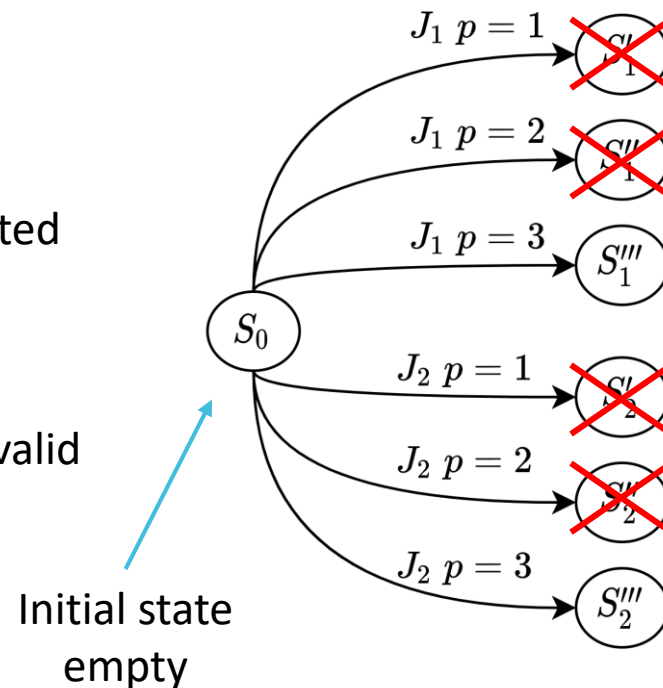


New SAG expansion rules

- Purge states by checking that
 - p cores available
 - More cores **not** available
 - Precedence constraints are respected
- Exploring more states

👍 Is **safe**, does not make analysis invalid

🗨 **Slower** and more **pessimistic**



New SAG merge rules

New SAG merge rules

- We have to merge

New SAG merge rules

- We have to merge
 - Availability times


New SAG merge rules

- We have to merge
 - Availability times
 - Groups of cores


New SAG merge rules

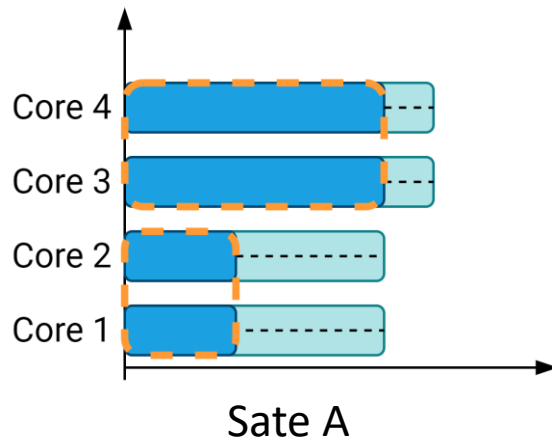
- We have to merge
 - Availability times
 - Groups of cores
 - Certainly running jobs

New SAG merge rules

- We have to merge
 - Availability times  Extend the intervals
 - Groups of cores
 - Certainly running jobs

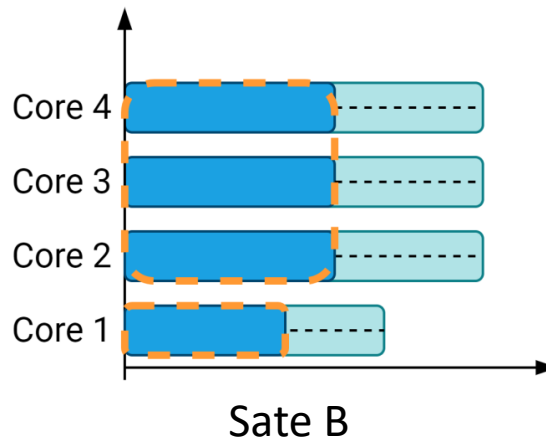
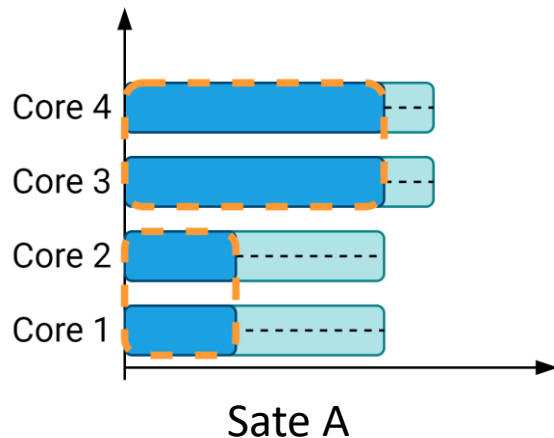
New SAG merge rules

- We have to merge
 - Availability times  Extend the intervals
 - Groups of cores
 - Certainly running jobs




New SAG merge rules

- We have to merge
 - Availability times → Extend the intervals
 - Groups of cores
 - Certainly running jobs



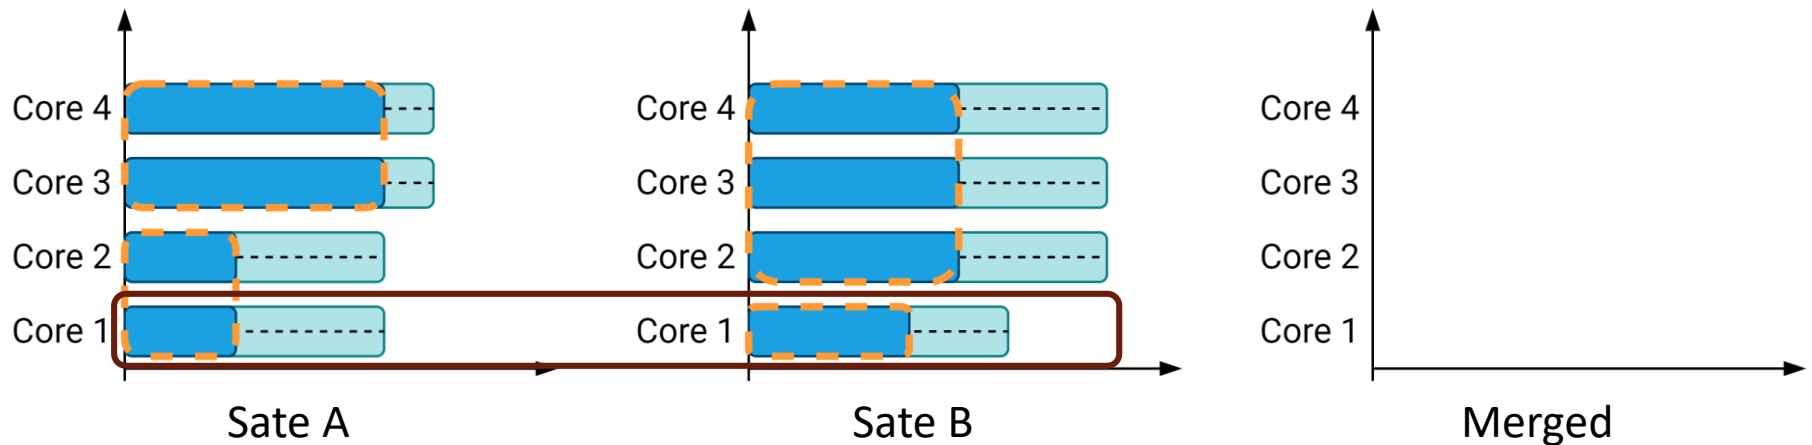
New SAG merge rules

- We have to merge
 - Availability times → Extend the intervals
 - Groups of cores
 - Certainly running jobs





New SAG merge rules

- We have to merge
 - Availability times \longrightarrow Extend the intervals
 - Groups of cores
 - Certainly running jobs




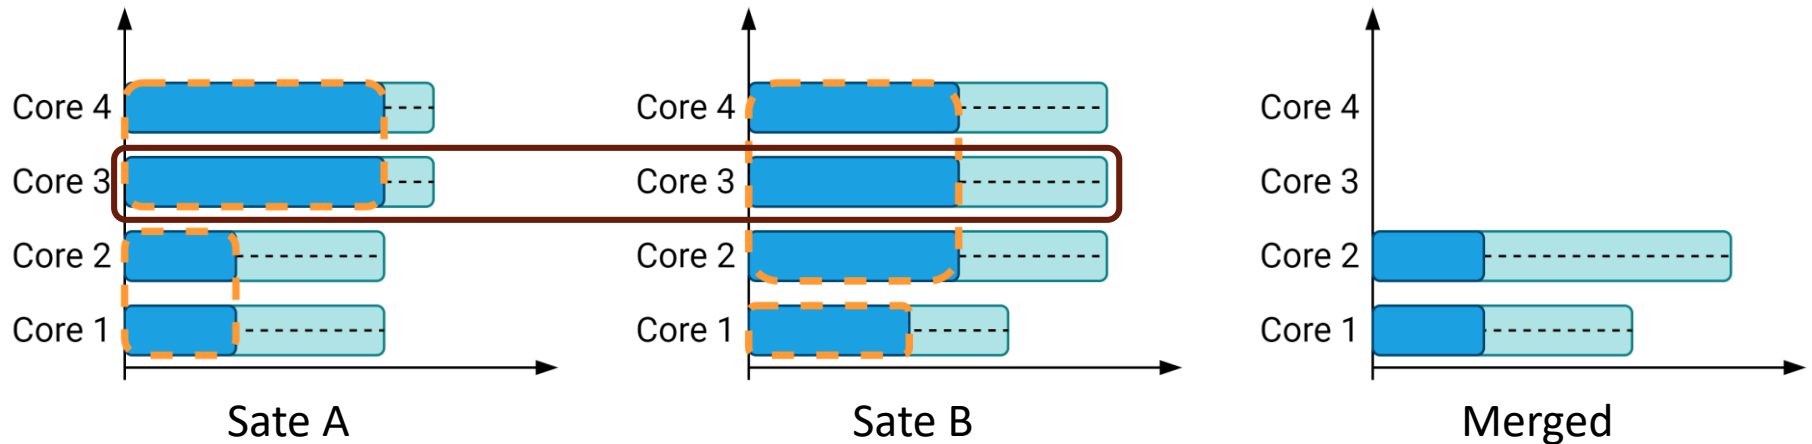
New SAG merge rules

- We have to merge
 - Availability times  Extend the intervals
 - Groups of cores
 - Certainly running jobs




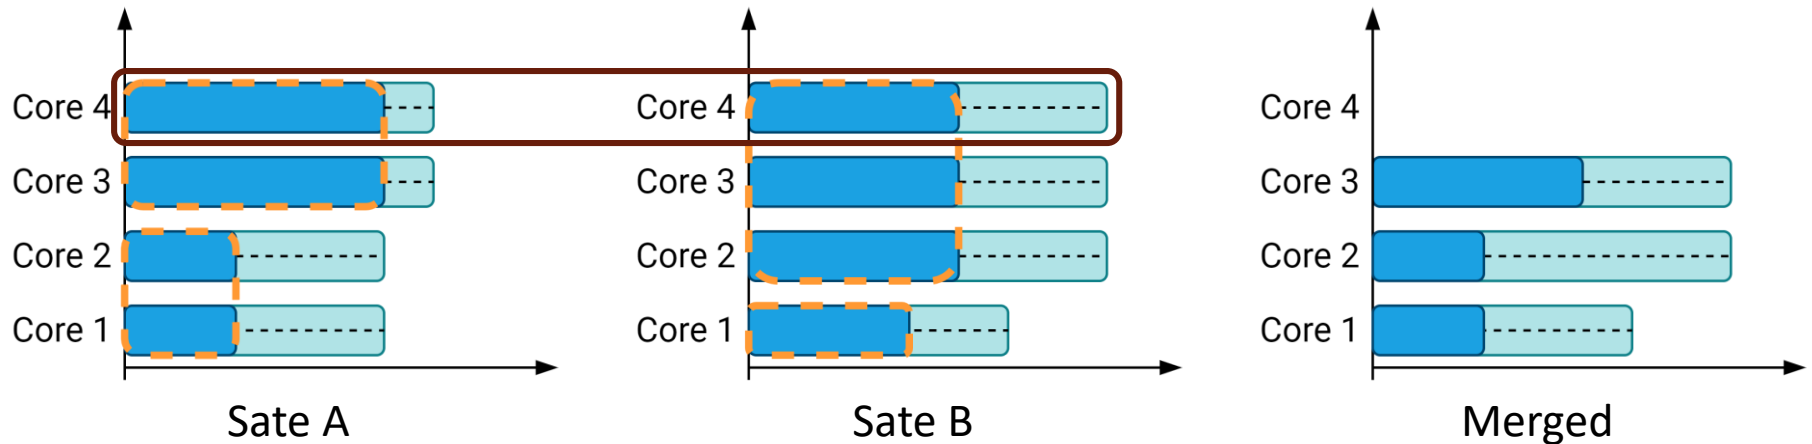
New SAG merge rules

- We have to merge
 - Availability times  Extend the intervals
 - Groups of cores
 - Certainly running jobs




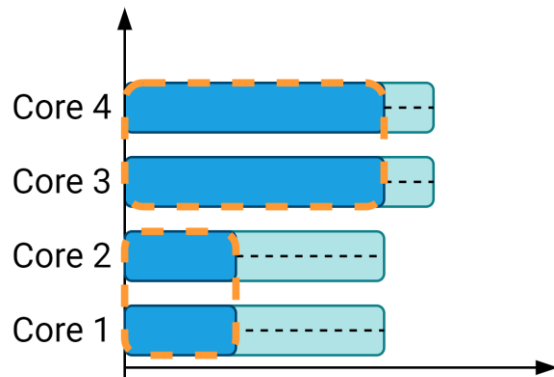
New SAG merge rules

- We have to merge
 - Availability times  Extend the intervals
 - Groups of cores
 - Certainly running jobs

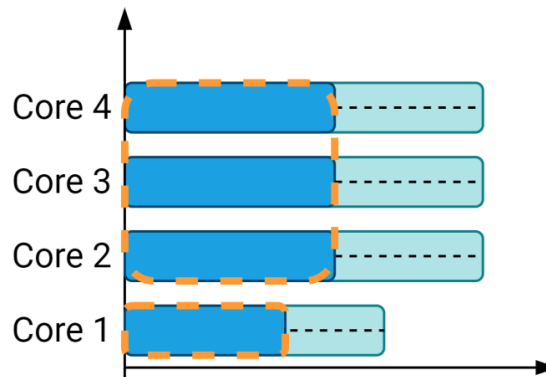


New SAG merge rules

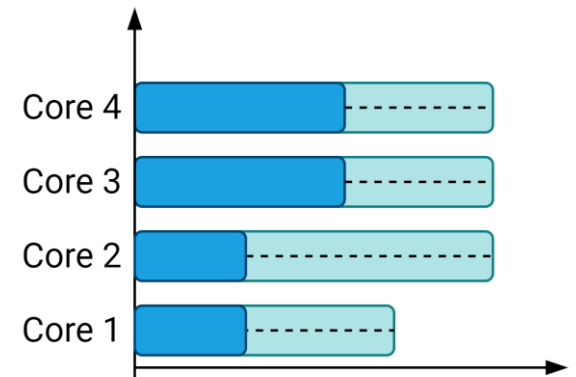
- We have to merge
 - Availability times  Extend the intervals
 - Groups of cores
 - Certainly running jobs



State A



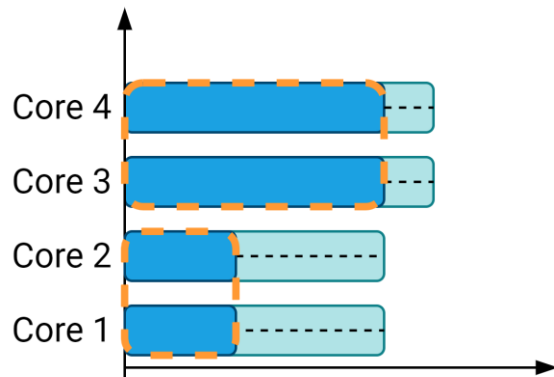
State B



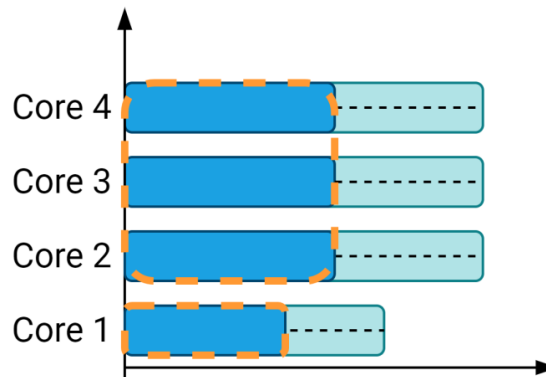
Merged

New SAG merge rules

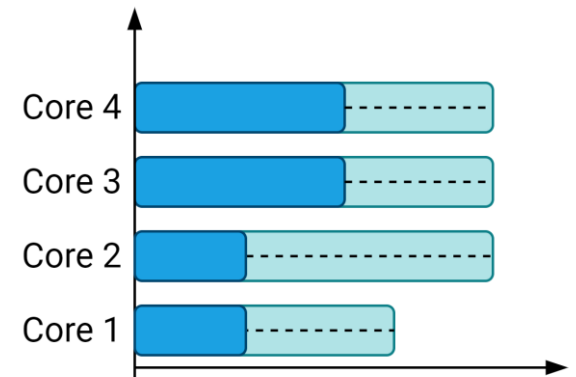
- We have to merge
 - Availability times \longrightarrow Extend the intervals
 - Groups of cores \longrightarrow Break the groups into same size
 - Certainly running jobs



State A



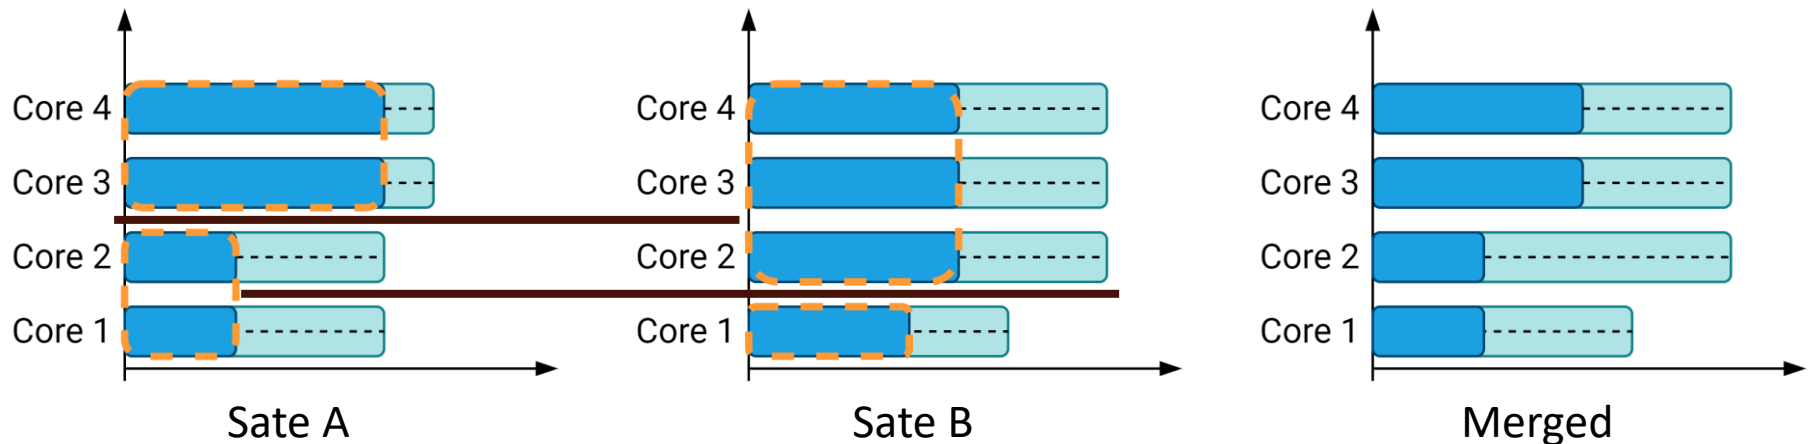
State B



Merged


New SAG merge rules

- We have to merge
 - Availability times \longrightarrow Extend the intervals
 - Groups of cores \longrightarrow Break the groups into same size
 - Certainly running jobs

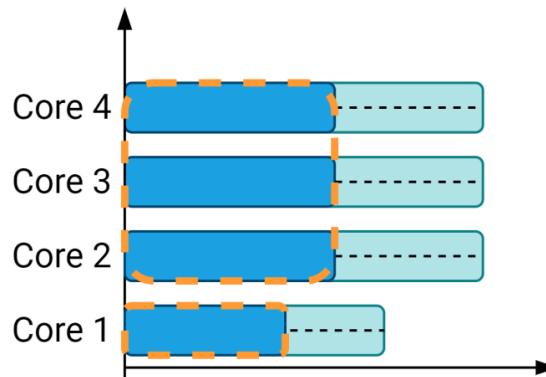


New SAG merge rules

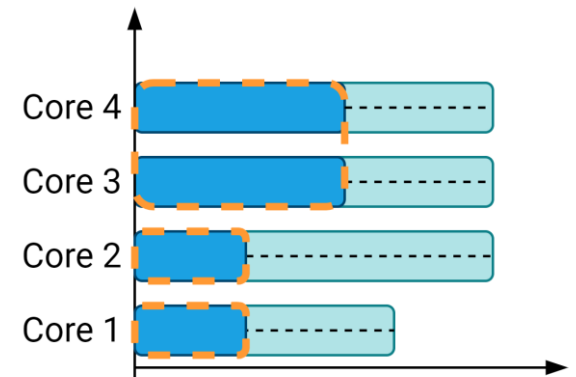
- We have to merge
 - Availability times \longrightarrow Extend the intervals
 - Groups of cores \longrightarrow Break the groups into same size
 - Certainly running jobs



State A



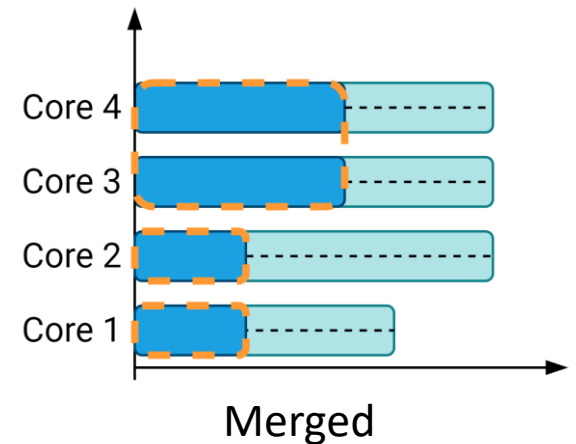
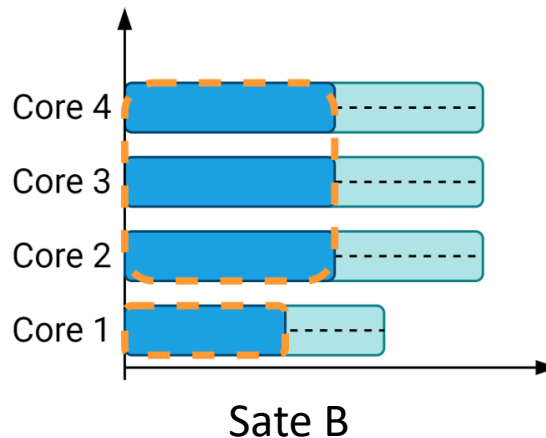
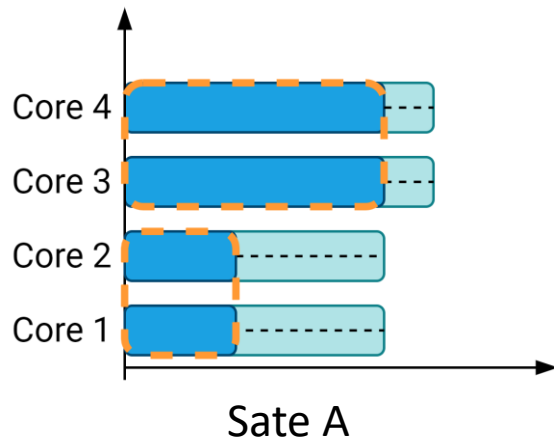
State B



Merged

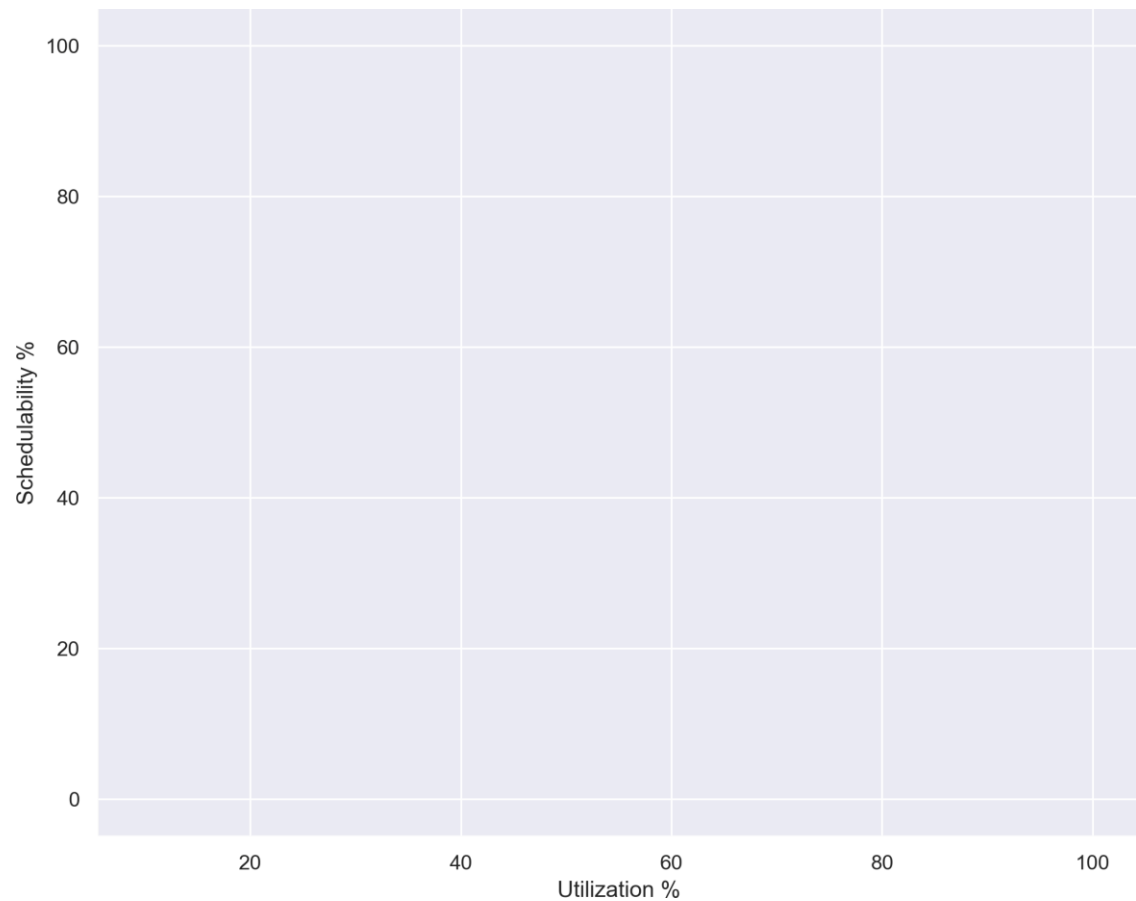
New SAG merge rules

- We have to merge
 - Availability times \longrightarrow Extend the intervals
 - Groups of cores \longrightarrow Break the groups into same size
 - Certainly running jobs \longrightarrow Keep only jobs running in both states (intersect)



Our analysis results

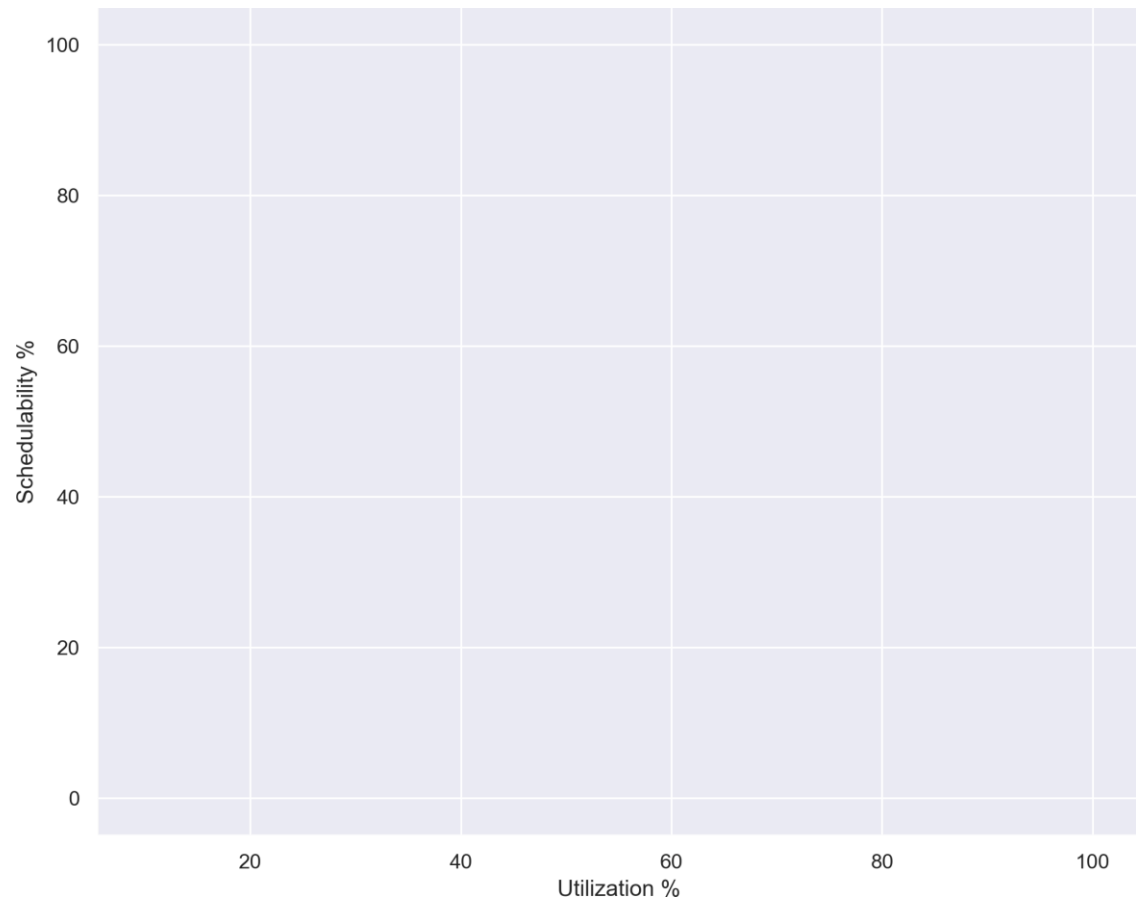
Our analysis results



Our analysis results

Randomly generated task sets

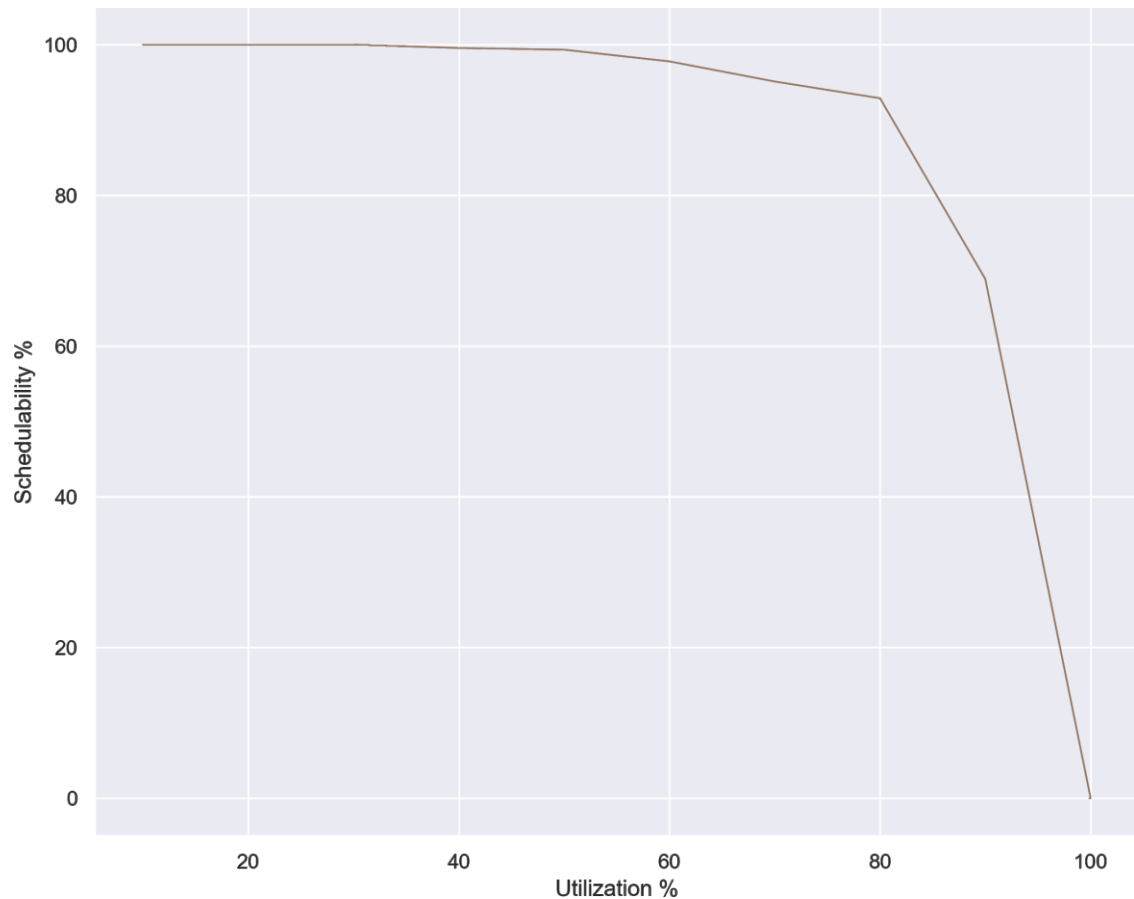
- System processors: 8
- System tasks: 20 rigid tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets


- System processors: 8
- System tasks: 20 rigid tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets

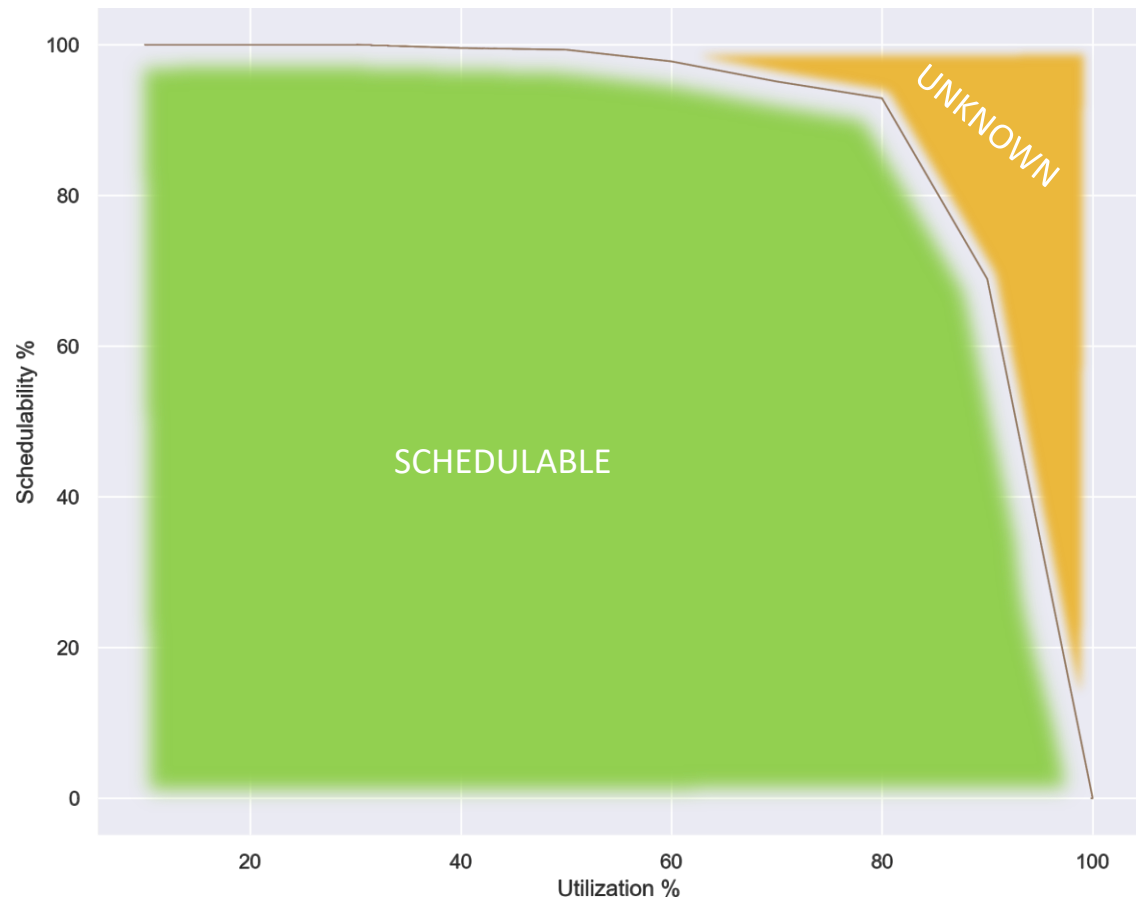
- System processors: 8
- System tasks: 20 rigid tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets

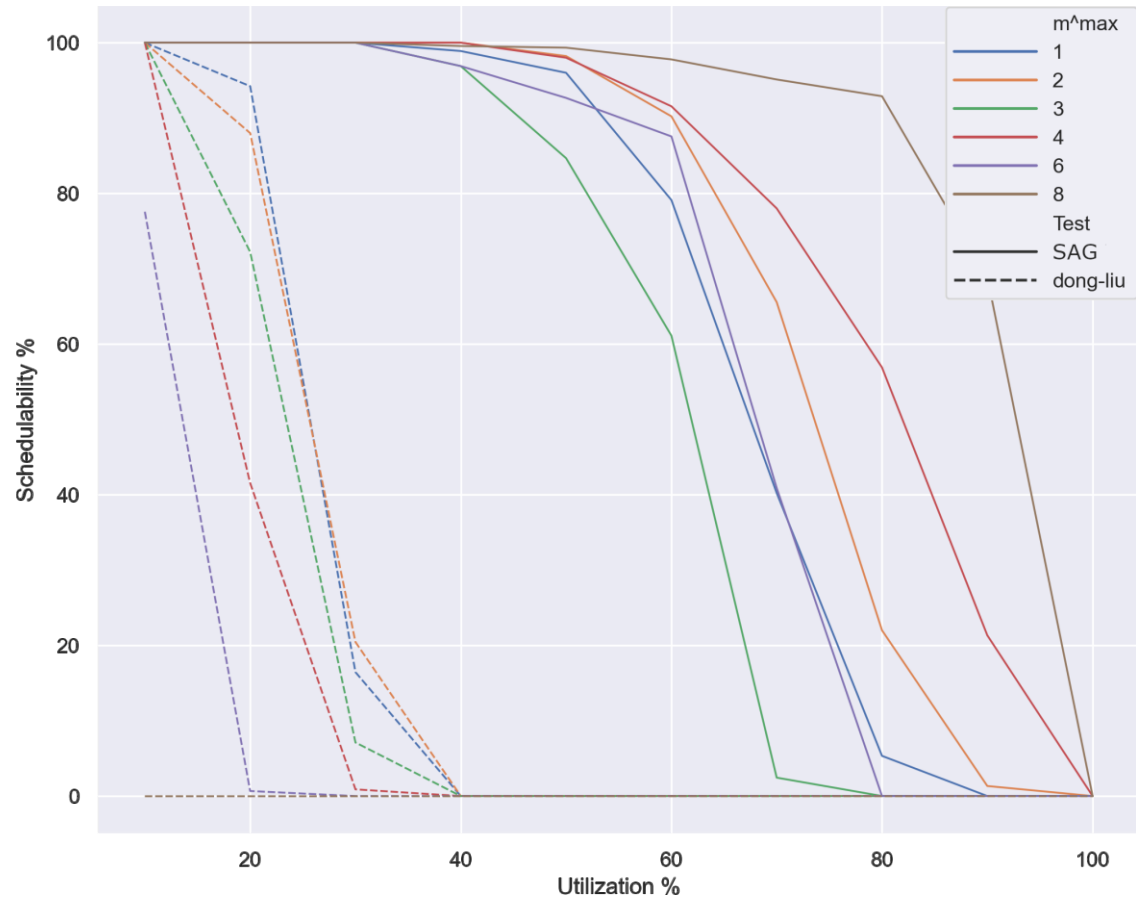
- System processors: 8
- System tasks: 20 rigid tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets

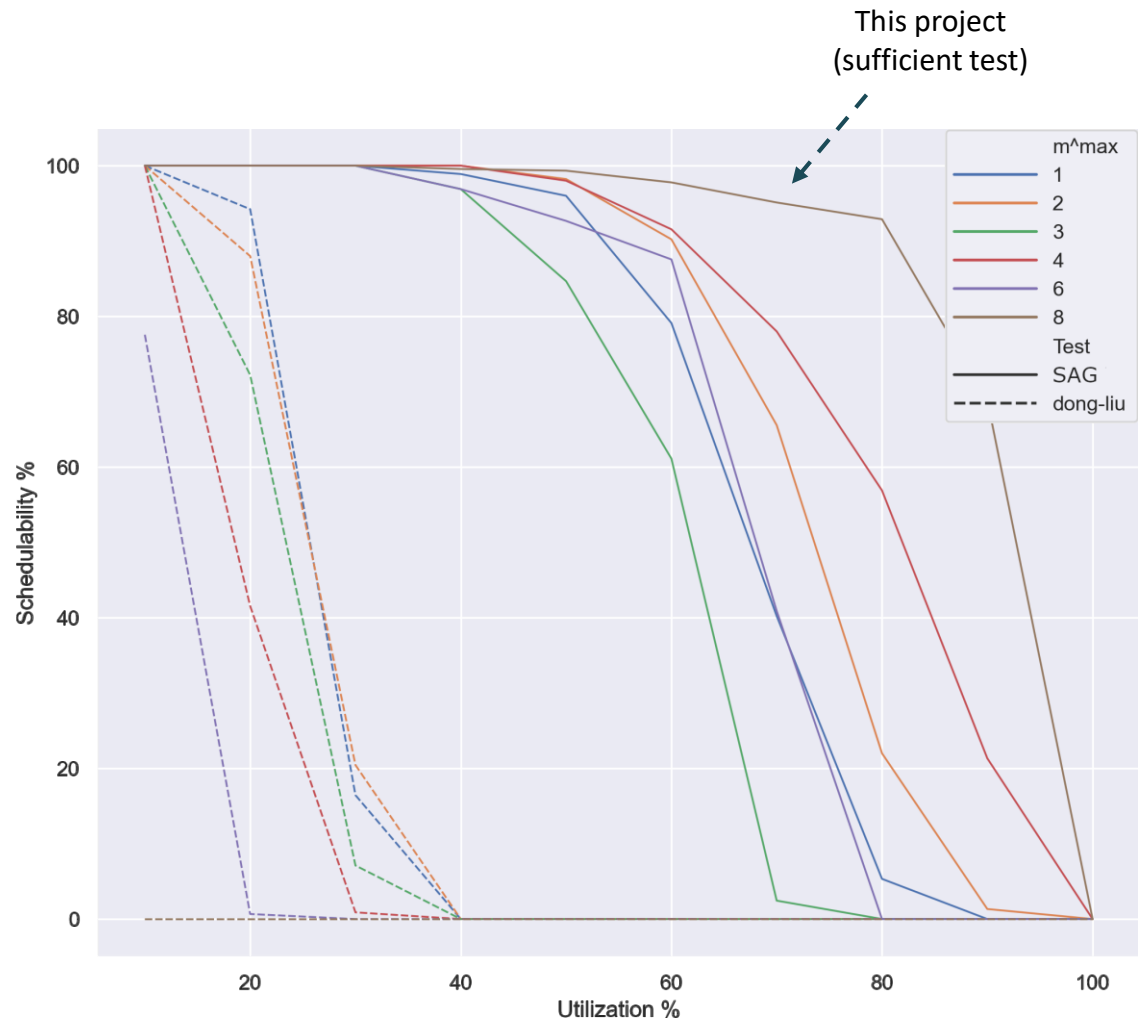
- System processors: 8
- System tasks: 20 rigid tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets

- System processors: 8
- System tasks: 20 rigid tasks
- Execution time variation: 50%

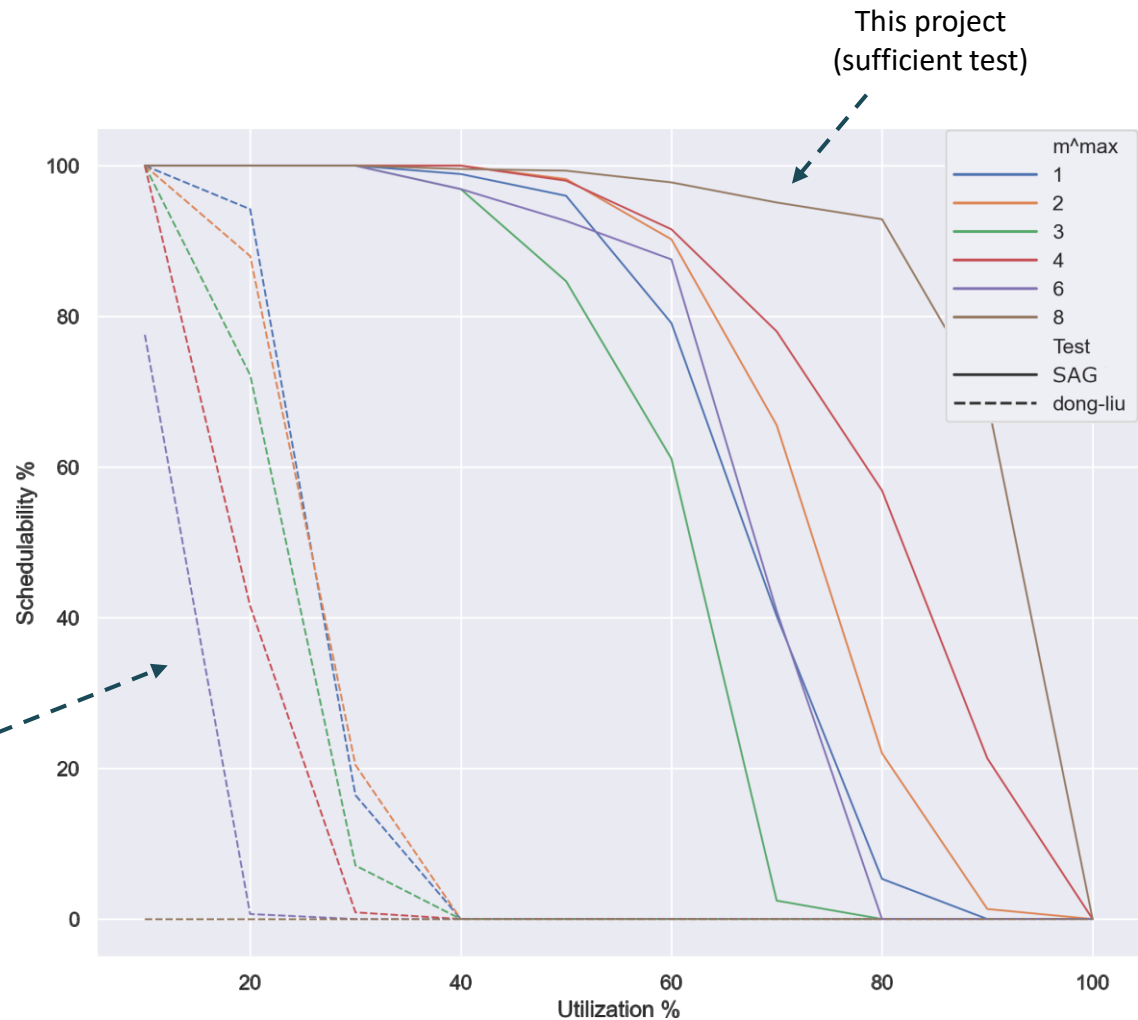


Our analysis results

Randomly generated task sets

- System processors: 8
- System tasks: 20 rigid tasks
- Execution time variation: 50%


State-of-the-art
(sufficient test)



Our analysis results

Randomly generated task sets


- System processors: 8
- System tasks: 20 **modal** tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets


- System processors: 8
- System tasks: 20 **modal** tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets


- System processors: 8
- System tasks: 20 **modal** tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets


- System processors: 8
- System tasks: 20 **modal** tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets


- System processors: 8
- System tasks: 20 **modal** tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets


- System processors: 8
- System tasks: 20 **modal** tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets


- System processors: 8
- System tasks: 20 **modal** tasks
- Execution time variation: 50%



Our analysis results

Randomly generated task sets

- System processors: 4
- System tasks: 4 moldable tasks
- Execution time variation: 50%
- System utilization: 70%




Agenda

- ~~Gang schedulability analysis~~
- New scheduling policy


JLFP limitations with moldable gang

JLFP limitations with moldable gang




	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

JLFP limitations with moldable gang




	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

JLFP limitations with moldable gang




	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

JLFP limitations with moldable gang




	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

JLFP limitations with moldable gang




	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

JLFP limitations with moldable gang




	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

JLFP limitations with moldable gang



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20


JLFP limitations with moldable gang



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler


- Reservation-based



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler


- Reservation-based



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler


- Reservation-based



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler


- Reservation-based
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler


- Reservation-based
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler


- Reservation-based
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler


- Reservation-based
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler

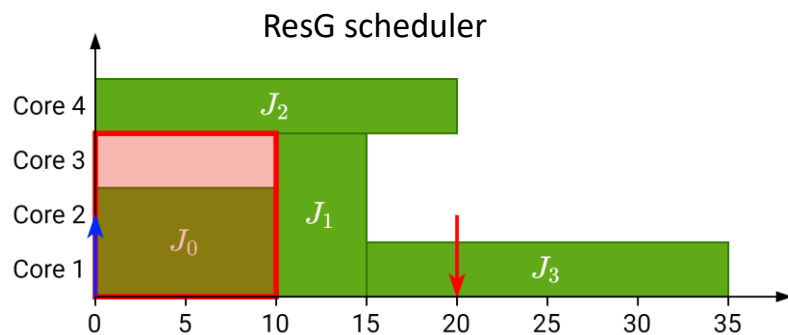
- Reservation-based
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Reservation-based gang scheduler

- Reservation-based
- Reserve cores of higher-priority tasks and distribute the remaining ones among lower priority tasks
- Non-work conserving scheduler



	Priority	Cores	Deadline	Execution time
J_0	High	2	∞	10
J_1	Mid-high	3	20	5
J_2	Mid-low	1	∞	20
J_3	Low	1	∞	20

Results JLFP vs ResG in simulator

Results JLFP vs ResG in simulator


- Evaluated in simulator

Results JLFP vs ResG in simulator

- Evaluated in simulator
- Randomly generated task sets


Results JLFP vs ResG in simulator

- Evaluated in simulator
- Randomly generated task sets




Results JLFP vs ResG in simulator

- Evaluated in simulator
- Randomly generated task sets




Results JLFP vs ResG in simulator

- Evaluated in simulator
- Randomly generated task sets



Results JLFP vs ResG in simulator

- Evaluated in simulator
- Randomly generated task sets



Conclusions

Conclusions

- With a better scheduling policy one can improve the schedulability of moldable gang tasks

Summary

Summary

- A new analysis for gang tasks using SAG has been defined

Summary

- A new analysis for gang tasks using SAG has been defined
- A new scheduling policy that uses gang moldable properties has been created

Next steps

Next steps

- Further reduce sources of pessimism



Next steps

- Further reduce sources of pessimism
- Provide analysis for ResG scheduler and respective proofs

Next steps

- Further reduce sources of pessimism
- Provide analysis for ResG scheduler and respective proofs
- Thorough evaluation of results using SURFSara cluster


Questions?



How does SAG work?


How does SAG work?

J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	100	2



How does SAG work?


J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	100	2



[5, 15]
[10, 20]

How does SAG work?


J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	100	2



[5, 15]
[10, 20]

How does SAG work?



J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	100	2



[5, 15]
[10, 20]


How does SAG work?

J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	00	2




How does SAG work?

J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	00	2





How does SAG work?

J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	00	2





How does SAG work?

J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	00	2



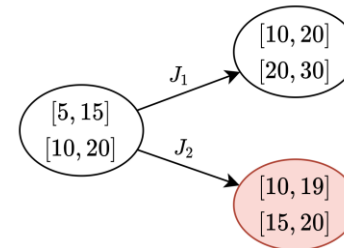

How does SAG work?

J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	00	2





How does SAG work?

J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	00	2



How does SAG work?

J_i	C_i^{min}	C_i^{max}	r_i	d_i	P_i
J_1	10	15	10	30	1
J_2	5	5	0	00	2



SAG analysis changes for gang

SAG analysis changes for gang

$$EST_i = \max\{R_i^{\min}, A_1^{\min}\}$$

$$LST_i = \min\{t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \max\{A_1^{\max}, \min\{R_x^{\max} \mid J_x \in \mathcal{R}^p\}\}$$

$$t_{high} = \min\{th_x(J_i) \mid J_x \in \mathcal{R}^p \wedge p_x < p_i\}$$

$$th_x(J_i) = \max\{r_x^{\max}, \max\{LFT_y^* \mid J_y \in pred(J_x) \setminus pred(J_i)\}\}$$

SAG analysis changes for gang

$$EST_i = \max\{R_i^{\min}, A_1^{\min}\}$$

$$LST_i = \min\{t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \max\{A_1^{\max}, \min\{R_x^{\max} \mid J_x \in \mathcal{R}^p\}\}$$

$$t_{high} = \min_{\infty}\{th_x(J_i) \mid J_x \in \mathcal{R}^p \wedge p_x < p_i\}$$

$$th_x(J_i) = \max\{r_x^{\max}, \max_{\infty}\{LFT_y^* \mid J_y \in pred(J_x) \setminus pred(J_i)\}\}$$

$$EST_i^p = \max\{R_i^{\min}, t_{gang}\}$$

$$LST_i^p = \min\{t_{avail}, t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \min_{J_j \in \mathcal{R}^v} \left\{ \max\{R_j^{\max}, A_{m_j^{\min}}^{\max}\} \right\}$$

$$t_{high} = \min_{J_j \in \{hp_i \cap \mathcal{R}^v\}} \left\{ th_x(J_i, J_j), \max_{\infty}\{LFT_y^* \mid J_y \in pred(J_j) \setminus pred(J_i)\} \right\}$$

$$th(J_i, J_j) = \begin{cases} r_j^{\max} & \text{if } m_j^{\min} \leq p \\ \max\{r_j^{\max}, A_{m_j^{\min}}^{\max}\} & \text{otherwise} \end{cases}$$

$$t_{gang} = \begin{cases} A_p^{\min} & \text{if } p = m_i^{\max} \\ A_p^{\text{exact}} & \text{otherwise} \end{cases} \quad t_{avail} = \begin{cases} A_{p+1}^{\max} - 1 & \text{if } p < m_i^{\max} \\ +\infty & \text{otherwise} \end{cases}$$

SAG analysis changes for gang

$$EST_i = \max\{R_i^{\min}, A_1^{\min}\}$$

$$LST_i = \min\{t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \max\{A_1^{\max}, \min\{R_x^{\max} \mid J_x \in \mathcal{R}^p\}\}$$

$$t_{high} = \min_{\infty}\{th_x(J_i) \mid J_x \in \mathcal{R}^p \wedge p_x < p_i\}$$

$$th_x(J_i) = \max\left\{r_x^{\max}, \max_{\infty}\{LFT_y^* \mid J_y \in pred(J_x) \setminus pred(J_i)\}\right\}$$

$$EST_i^p = \max\{R_i^{\min}, t_{gang}\}$$

$$LST_i^p = \min\{t_{avail}, t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \min_{J_j \in \mathcal{R}^p} \left\{ \max\left\{R_j^{\max}, A_{m_j^{\min}}^{\max}\right\} \right\}$$

$$t_{high} = \min_{J_j \in \{hp_i \cap \mathcal{R}^p\}} \left\{ th_x(J_i, J_j), \max_{\infty}\{LFT_y^* \mid J_y \in pred(J_j) \setminus pred(J_i)\} \right\}$$

$$th(J_i, J_j) = \begin{cases} r_j^{\max} & \text{if } m_j^{\min} \leq p \\ \max\left\{r_j^{\max}, A_{m_j^{\min}}^{\max}\right\} & \text{otherwise} \end{cases}$$

$$t_{gang} = \begin{cases} A_p^{\min} & \text{if } p = m_i^{\max} \\ A_p^{\text{exact}} & \text{otherwise} \end{cases} \quad t_{avail} = \begin{cases} A_{p+1}^{\max} - 1 & \text{if } p < m_i^{\max} \\ +\infty & \text{otherwise} \end{cases}$$

Check if execution with p cores is possible

SAG analysis changes for gang

$$EST_i = \max\{R_i^{\min}, A_1^{\min}\}$$

$$LST_i = \min\{t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \max\{A_1^{\max}, \min\{R_x^{\max} \mid J_x \in \mathcal{R}^p\}\}$$

$$t_{high} = \min_{\infty}\{th_x(J_i) \mid J_x \in \mathcal{R}^p \wedge p_x < p_i\}$$

$$th_x(J_i) = \max\left\{r_x^{\max}, \max_{\infty}\{LFT_y^* \mid J_y \in pred(J_x) \setminus pred(J_i)\}\right\}$$

$$EST_i^p = \max\{R_i^{\min}, t_{gang}\}$$

$$LST_i^p = \min\{t_{avail}, t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \min_{J_j \in \mathcal{R}^p} \left\{ \max\left\{R_j^{\max}, A_{m_j^{\min}}^{\max}\right\} \right\}$$

$$t_{high} = \min_{J_j \in \{hp_i \cap \mathcal{R}^p\}} \left\{ th_x(J_i, J_j) \right. \\ \left. \max_{\infty}\{LFT_y^* \mid J_y \in pred(J_j) \setminus pred(J_i)\} \right\}$$

$$th(J_i, J_j) = \begin{cases} r_j^{\max} & \text{if } m_j^{\min} \leq p \\ \max\left\{r_j^{\max}, A_{m_j^{\min}}^{\max}\right\} & \text{otherwise} \end{cases}$$

$$t_{gang} = \begin{cases} A_p^{\min} & \text{if } p = m_i^{\max} \\ A_p^{\text{exact}} & \text{otherwise} \end{cases} \quad t_{avail} = \begin{cases} A_{p+1}^{\max} - 1 & \text{if } p < m_i^{\max} \\ +\infty & \text{otherwise} \end{cases}$$

p cores
available

SAG analysis changes for gang

$$EST_i = \max\{R_i^{\min}, A_1^{\min}\}$$

$$LST_i = \min\{t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \max\{A_1^{\max}, \min\{R_x^{\max} \mid J_x \in \mathcal{R}^p\}\}$$

$$t_{high} = \min_{\infty}\{th_x(J_i) \mid J_x \in \mathcal{R}^p \wedge p_x < p_i\}$$

$$th_x(J_i) = \max\left\{r_x^{\max}, \max_{\infty}\{LFT_y^* \mid J_y \in pred(J_x) \setminus pred(J_i)\}\right\}$$

$$EST_i^p = \max\{R_i^{\min}, t_{gang}\}$$

$$LST_i^p = \min\{t_{avail}, t_{wc}, t_{high} - 1\}$$

$$t_{wc} = \min_{J_j \in \mathcal{R}^p} \left\{ \max\left\{R_j^{\max}, A_{m_j^{\min}}^{\max}\right\} \right\}$$

$$t_{high} = \min_{J_j \in \{hp_i \cap \mathcal{R}^p\}} \left\{ th_x(J_i, J_j), \max_{\infty}\{LFT_y^* \mid J_y \in pred(J_j) \setminus pred(J_i)\} \right\}$$

$$t_h(J_i, J_j) = \begin{cases} r_j^{\max} & \text{if } m_j^{\min} \leq p \\ \max\{r_j^{\max}, A_{m_j^{\min}}^{\max}\} & \text{otherwise} \end{cases}$$

$$t_{gang} = \begin{cases} A_p^{\min} & \text{if } p = m_i^{\max} \\ A_p^{\text{exact}} & \text{otherwise} \end{cases}$$

$$t_{avail} = \begin{cases} A_{p+1}^{\max} - 1 & \text{if } p < m_i^{\max} \\ +\infty & \text{otherwise} \end{cases}$$

$p + 1$ cores **not** available